

# Linux的调度分析

## 一、就绪队列

内核为每个CPU创建一个进程就绪队列，该队列上的进程均有该CPU执行。

per-cpu变量在每个CPU上都有一个副本，对它的访问几乎不需要锁，因为每个CPU都在自己的副本上工作。

```
1  /*在 kernel/sched/core.c中定义*/
2  DEFINE_PER_CPU_SHARED_ALIGNED(struct rq, runqueues);
3
4  /* 宏定义展开在include/linux/percpu-defs.h
5   * 静态分配per_cpu数组，数组名为name，结构类型为type
6   */
7  #define DEFINE_PER_CPU_SHARED_ALIGNED(type, name) \
8      DEFINE_PER_CPU_SECTION(type, name, PER_CPU_SHARED_ALIGNED_SECTION) \
9      ____cacheline_aligned_in_smp
```

每个数组元素就是一个就绪队列，对应一个cpu。

```
1  /*
2   * This is the main, per-CPU runqueue data structure.
3   *
4   * Locking rule: those places that want to lock multiple runqueues
5   * (such as the load balancing or the thread migration code), lock
6   * acquire operations must be ordered by ascending &runqueue.
7   */
8  struct rq {
9      /* runqueue lock: */
10     raw_spinlock_t    lock;
11
12     /*
13      * nr_running and cpu_load should be in the same cacheline because
14      * remote CPUs use both these fields when doing load calculation.
15      */
16     unsigned int      nr_running;
17 #ifdef CONFIG_NUMA_BALANCING
18     unsigned int      nr_numa_running;
19     unsigned int      nr_preferred_running;
20     unsigned int      numa_migrate_on;
21 #endif
22 #ifdef CONFIG_NO_HZ_COMMON
23 #ifdef CONFIG_SMP
24     unsigned long      last_blocked_load_update_tick;
25     unsigned int      has_blocked_load;
26     call_single_data_t nohz_csd;
27 #endif /* CONFIG_SMP */
28     unsigned int      nohz_tick_stopped;
29     atomic_t          nohz_flags;
30 #endif /* CONFIG_NO_HZ_COMMON */
31 }
```

```

32 #ifdef CONFIG_SMP
33     unsigned int        ttwu_pending;
34 #endif
35     u64                nr_switches;
36
37 #ifdef CONFIG_UCLAMP_TASK
38     /* Utilization clamp values based on CPU's RUNNABLE tasks */
39     struct uclamp_rq    uclamp[UCLAMP_CNT] ____cacheline_aligned;
40     unsigned int        uclamp_flags;
41 #define UCLAMP_FLAG_IDLE 0x01
42 #endif
43
44     struct cfs_rq        cfs;                /* 嵌入普通进程的cfs调度队列 */
45     struct rt_rq         rt;                /* 实时进程的实时调度策略调度队列 */
46     struct dl_rq         dl;                /* 空闲进程的调度队列 */
47
48 #ifdef CONFIG_FAIR_GROUP_SCHED
49     /* list of leaf cfs_rq on this CPU: */
50     struct list_head     leaf_cfs_rq_list;
51     struct list_head     *tmp_alone_branch;
52 #endif /* CONFIG_FAIR_GROUP_SCHED */
53
54     /*
55      * This is part of a global counter where only the total sum
56      * over all CPUs matters. A task can increase this counter on
57      * one CPU and if it got migrated afterwards it may decrease
58      * it on another CPU. Always updated under the runqueue lock:
59      */
60     unsigned long        nr_uninterruptible;
61
62     struct task_struct    __rcu    *curr;
63     struct task_struct    *idle;
64     struct task_struct    *stop;
65     unsigned long        next_balance;
66     struct mm_struct      *prev_mm;
67
68     unsigned int        clock_update_flags;
69     u64                clock;
70     /* Ensure that all clocks are in the same cache line */
71     u64                clock_task ____cacheline_aligned;
72     u64                clock_pelt;
73     unsigned long        lost_idle_time;
74
75     atomic_t            nr_iowait;
76
77 #ifdef CONFIG_MEMBARRIER
78     int membarrier_state;
79 #endif
80
81 #ifdef CONFIG_SMP
82     struct root_domain    *rd;
83     struct sched_domain    __rcu    *sd;
84
85     unsigned long        cpu_capacity;
86     unsigned long        cpu_capacity_orig;
87
88     struct callback_head    *balance_callback;
89

```

```

90     unsigned char    nohz_idle_balance;
91     unsigned char    idle_balance;
92
93     unsigned long     misfit_task_load;
94
95     /* For active balancing */
96     int               active_balance;
97     int               push_cpu;
98     struct cpu_stop_work  active_balance_work;
99
100    /* CPU of this runqueue: */
101    int                cpu;
102    int                online;
103
104    struct list_head cfs_tasks;
105
106    struct sched_avg    avg_rt;
107    struct sched_avg    avg_dl;
108    #ifdef CONFIG_HAVE_SCHED_AVG_IRQ
109    struct sched_avg    avg_irq;
110    #endif
111    #ifdef CONFIG_SCHED_THERMAL_PRESSURE
112    struct sched_avg    avg_thermal;
113    #endif
114    u64                 idle_stamp;
115    u64                 avg_idle;
116
117    /* This is used to determine avg_idle's max value */
118    u64                 max_idle_balance_cost;
119    #endif /* CONFIG_SMP */
120
121    #ifdef CONFIG_IRQ_TIME_ACCOUNTING
122    u64                 prev_irq_time;
123    #endif
124    #ifdef CONFIG_PARAVIRT
125    u64                 prev_steal_time;
126    #endif
127    #ifdef CONFIG_PARAVIRT_TIME_ACCOUNTING
128    u64                 prev_steal_time_rq;
129    #endif
130
131    /* calc_load related fields */
132    unsigned long        calc_load_update;
133    long                 calc_load_active;
134
135    #ifdef CONFIG_SCHED_HRTICK
136    #ifdef CONFIG_SMP
137    call_single_data_t   hrtick_csd;
138    #endif
139    struct hrtimer        hrtick_timer;
140    #endif
141
142    #ifdef CONFIG_SCHEDSTATS
143    /* latency stats */
144    struct sched_info    rq_sched_info;
145    unsigned long long    rq_cpu_time;
146    /* could above be rq->cfs_rq.exec_clock + rq->rt_rq.rt_runtime ? */
147

```

```

148     /* sys_sched_yield() stats */
149     unsigned int      yld_count;
150
151     /* schedule() stats */
152     unsigned int      sched_count;
153     unsigned int      sched_goidle;
154
155     /* try_to_wake_up() stats */
156     unsigned int      ttwu_count;
157     unsigned int      ttwu_local;
158 #endif
159
160 #ifdef CONFIG_CPU_IDLE
161     /* Must be inspected within a rcu lock section */
162     struct cpuidle_state *idle_state;
163 #endif
164 }

```

## 1.1 普通进程cfs就绪队列

`cfs_rq`和`rt_rq`以及`dl_rq`都定义在 `kernel/sched/sched.h` 中。

```

1  /* CFS-related fields in a runqueue */
2  struct cfs_rq {
3      struct load_weight  load;                /* load维护了所有这些进程的
累积负荷值 */
4      unsigned int        nr_running;          /* nr_running计算了队列上
可运行进程的数目 */
5      unsigned int        h_nr_running;        /*
SCHED_{NORMAL,BATCH,IDLE} */
6      unsigned int        idle_h_nr_running;   /* SCHED_IDLE */
7
8      u64                  exec_clock;
9      u64                  min_vruntime;       /* 跟踪记录队列上所有进程的
最小虚拟运行时间 */
10 #ifndef CONFIG_64BIT
11     u64                  min_vruntime_copy;
12 #endif
13
14     struct rb_root_cached tasks_timeline;     /* 红黑树根以及待被调用的进
程所在的树节点 */
15
16     /*
17      * 'curr' points to currently running entity on this cfs_rq.
18      * It is set to NULL otherwise (i.e when none are currently running).
19      */
20     struct sched_entity *curr;               /* curr指向当前正运行的实体
*/
21     struct sched_entity *next;               /* next指向将被唤醒的进程
*/
22     struct sched_entity *last;               /* last指向唤醒next进程的进
程 */
23     struct sched_entity *skip;
24
25 #ifdef CONFIG_SCHED_DEBUG

```

```

26     unsigned int        nr_spread_over;
27 #endif
28
29 #ifdef CONFIG_SMP
30     /*
31      * CFS load tracking
32      */
33     struct sched_avg    avg;
34 #ifndef CONFIG_64BIT
35     u64        load_last_update_time_copy;
36 #endif
37     struct {
38         raw_spinlock_t    lock ____cacheline_aligned;
39         int                nr;
40         unsigned long      load_avg;
41         unsigned long      util_avg;
42         unsigned long      runnable_avg;
43     } removed;
44
45 #ifdef CONFIG_FAIR_GROUP_SCHED
46     unsigned long          tg_load_avg_contrib;
47     long                  propagate;
48     long                  prop_runnable_sum;
49
50     /*
51      *   h_load = weight * f(tg)
52      *
53      * where f(tg) is the recursive weight fraction assigned to
54      * this group.
55      */
56     unsigned long          h_load;
57     u64                    last_h_load_update;
58     struct sched_entity    *h_load_next;
59 #endif /* CONFIG_FAIR_GROUP_SCHED */
60 #endif /* CONFIG_SMP */
61
62 #ifdef CONFIG_FAIR_GROUP_SCHED
63     struct rq              *rq;    /* CPU runqueue to which this cfs_rq is attached
64     */
65
66     /*
67      * leaf cfs_rqs are those that hold tasks (lowest schedulable entity in
68      * a hierarchy). Non-leaf lrs hold other higher schedulable entities
69      * (like users, containers etc.)
70      *
71      * leaf_cfs_rq_list ties together list of leaf cfs_rq's in a CPU.
72      * This list is used during load balance.
73      */
74     int                    on_list;
75     struct list_head       leaf_cfs_rq_list;
76     struct task_group       *tg;    /* group that "owns" this runqueue */
77
78 #ifdef CONFIG_CFS_BANDWIDTH
79     int                    runtime_enabled;
80     s64                    runtime_remaining;
81
82     u64                    throttled_clock;
83     u64                    throttled_clock_task;

```

```

83     u64          throttled_clock_task_time;
84     int          throttled;
85     int          throttle_count;
86     struct list_head throttled_list;
87 #endif /* CONFIG_CFS_BANDWIDTH */
88 #endif /* CONFIG_FAIR_GROUP_SCHED */
89 };

```

其中 `rb_root_cached` 包含了两个成员：

一个是红黑树根 `rb_root`；

另一个是 `rb_node`，其总是设置为指向树最左边的结点，即最需要被调度的进程；

```

1  /*
2   * Leftmost-cached rbtrees.
3   *
4   * We do not cache the rightmost node based on footprint
5   * size vs number of potential users that could benefit
6   * from O(1) rb_last(). Just not worth it, users that want
7   * this feature can always implement the logic explicitly.
8   * Furthermore, users that want to cache both pointers may
9   * find it a bit asymmetric, but that's ok.
10  */
11  struct rb_root_cached {
12      struct rb_root rb_root;
13      struct rb_node *rb_leftmost;
14  };

```

## 1.2 实时进程rt就绪队列

```

1  /* Real-Time classes' related field in a runqueue: */
2  struct rt_rq {
3      struct rt_prio_array  active;
4      unsigned int          rt_nr_running;
5      unsigned int          rr_nr_running;
6  #if defined CONFIG_SMP || defined CONFIG_RT_GROUP_SCHED
7      struct {
8          int curr; /* highest queued rt task prio */
9  #ifdef CONFIG_SMP
10         int next; /* next highest */
11     #endif
12     } highest_prio;
13 #endif
14 #ifdef CONFIG_SMP
15     unsigned long rt_nr_migratory;
16     unsigned long rt_nr_total;
17     int overloaded;
18     struct plist_head pushable_tasks;
19 }

```

```

20 #endif /* CONFIG_SMP */
21     int            rt_queued;
22
23     int            rt_throttled;
24     u64            rt_time;
25     u64            rt_runtime;
26     /* Nests inside the rq lock: */
27     raw_spinlock_t  rt_runtime_lock;
28
29 #ifdef CONFIG_RT_GROUP_SCHED
30     unsigned long   rt_nr_boosted;
31
32     struct rq        *rq;
33     struct task_group *tg;
34 #endif
35 };

```

## 二、调度实体

### 2.1 普通进程调度实体

由于调度器可以操作比进程更一般的实体，因此需要一个适当的数据结构来描述此类实体

```

1  struct sched_entity {
2      /* For load-balancing: */
3      struct load_weight  load;                /* 用于负载均衡 */
4      struct rb_node      run_node;            /* run_node是标准的树
   结点，使得实体可以在红黑树上排序 */
5      struct list_head    group_node;
6      unsigned int        on_rq;                /* on_rq表示该实体当前
   是否在就绪队列上接受调度 */
7
8      u64                  exec_start;
9      u64                  sum_exec_runtime;
10     u64                  vruntime;            /* 在进程执行期间虚拟时
   钟上流逝的时间数量由vruntime统计 */
11     u64                  prev_sum_exec_runtime;
12
13     u64                  nr_migrations;
14
15     struct sched_statistics statistics;
16
17 #ifdef CONFIG_FAIR_GROUP_SCHED
18     int                   depth;
19     struct sched_entity   *parent;
20     /* rq on which this entity is (to be) queued: */
21     struct cfs_rq         *cfs_rq;
22     /* rq "owned" by this entity/group: */
23     struct cfs_rq         *my_q;
24     /* cached value of my_q->h_nr_running */
25     unsigned long         runnable_weight;
26 #endif
27

```

```

28 #ifdef CONFIG_SMP
29     /*
30      * Per entity load average tracking.
31      *
32      * Put into separate cache line so it does not
33      * collide with read-mostly values above.
34      */
35     struct sched_avg      avg;
36 #endif
37 };

```

`struct sched_entity` 该结构体有两个作用：

(1) 包含有进程调度的信息（比如进程的运行时间，睡眠时间等等，调度程序参考这些信息决定是否调度进程）

(2) 使用该结构体来组织进程

- `struct load_weight load` 指定了权重，决定了各个实体占队列总负荷的比例。计算负荷权重是调度器的一项重任，因为CFS所需的虚拟时钟的速度最终依赖于负荷。
- `run_node` 是红黑树节点，因此 `struct sched_entity` 调度实体将被组织成红黑树的形式，同时意味着普通进程也被组织成红黑树的形式。
- 在进程运行时，我们需要记录消耗的CPU时间，以用于完全公平调度器。`sum_exec_runtime` 即用于该目的。跟踪运行时间是由 `update_curr` 不断累积完成的。调度器中许多地方都会调用该函数，例如，新进程加入就绪队列时，或者周期性调度器中。每次调用时，会计算当前时间和 `exec_start` 之间的差值，`exec_start` 则更新到当前时间。差值则被加到 `sum_exec_runtime`。
- 在进程被撤销CPU时，其当前 `sum_exec_runtime` 值保存到 `prev_exec_runtime`。此后，在进程抢占时又需要该数据。但请注意，在 `prev_exec_runtime` 中保存 `sum_exec_runtime` 的值，并不意味着重置 `sum_exec_runtime`！原值保存下来，而 `um_exec_runtime` 则持续单调增长

## 2.2 实时进程rt调度实体

用于组织实时进程的调度。

```

1  struct sched_rt_entity {
2      struct list_head      run_list;
3      unsigned long         timeout;
4      unsigned long         watchdog_stamp;
5      unsigned int          time_slice;
6      unsigned short        on_rq;
7      unsigned short        on_list;
8
9      struct sched_rt_entity *back;
10 #ifdef CONFIG_RT_GROUP_SCHED
11     struct sched_rt_entity *parent;
12     /* rq on which this entity is (to be) queued: */
13     struct rt_rq           *rt_rq;
14     /* rq "owned" by this entity/group: */
15     struct rt_rq           *my_q;
16 #endif
17 } __randomize_layout;

```

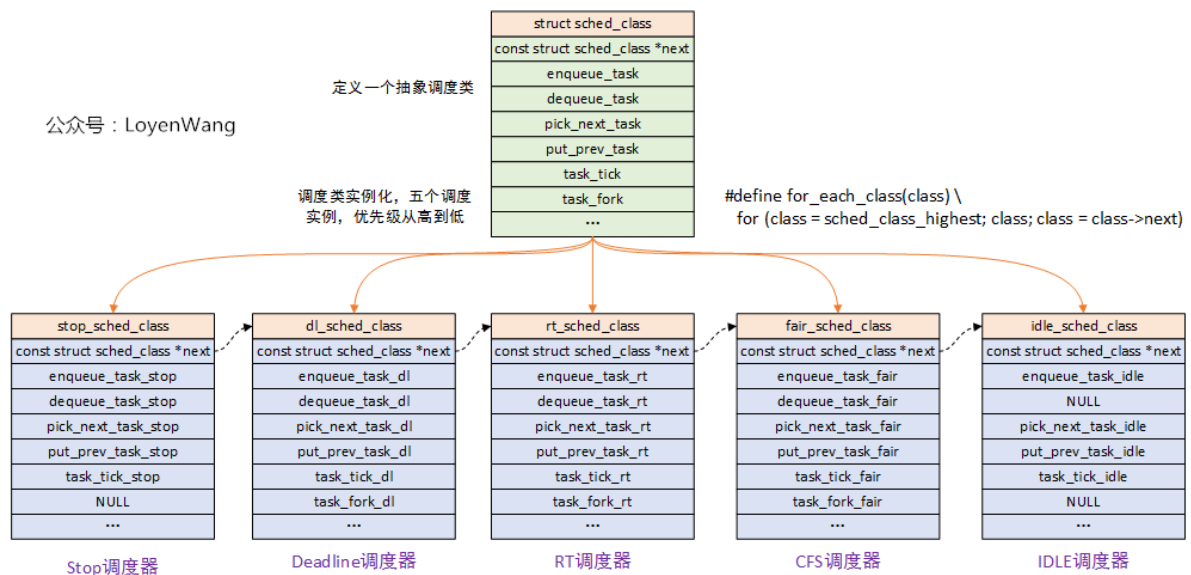
`struct list_head run_list` 表明 `rt_entity` 是由双链表管理，而不是红黑树。



## 三、调度类

### 3.1 调度类sched\_class

调度类整体关系：



kernel/sched/sched.h 中声明了调度类。

sched\_class 中定义了一堆函数指针，指针指向的函数就是调度策略的具体实现，所有和进程调度有关的函数都直接或者间接调用了这些成员函数，来实现进程调度。此外，每个进程描述符中都包含一个指向该结构体类型的指针sched\_class，指向了所采用的调度类。

```
1 struct sched_class {
2
3 #ifdef CONFIG_UCLAMP_TASK
4     int uclamp_enabled;
5 #endif
6
7     void (*enqueue_task) (struct rq *rq, struct task_struct *p, int flags);
8     void (*dequeue_task) (struct rq *rq, struct task_struct *p, int flags);
9     void (*yield_task) (struct rq *rq);
10    bool (*yield_to_task)(struct rq *rq, struct task_struct *p);
11
12    void (*check_preempt_curr)(struct rq *rq, struct task_struct *p, int
13 flags);
14
15    struct task_struct *(*pick_next_task)(struct rq *rq);
16
17    void (*put_prev_task)(struct rq *rq, struct task_struct *p);
18    void (*set_next_task)(struct rq *rq, struct task_struct *p, bool first);
19
20 #ifdef CONFIG_SMP
21     int (*balance)(struct rq *rq, struct task_struct *prev, struct rq_flags
22 *rf);
23     int (*select_task_rq)(struct task_struct *p, int task_cpu, int sd_flag,
24 int flags);
25     void (*migrate_task_rq)(struct task_struct *p, int new_cpu);
26     void (*task_woken)(struct rq *this_rq, struct task_struct *task);
27 }
```

```

25
26     void (*set_cpus_allowed)(struct task_struct *p,
27                             const struct cpumask *newmask);
28
29     void (*rq_online)(struct rq *rq);
30     void (*rq_offline)(struct rq *rq);
31 #endif
32
33     void (*task_tick)(struct rq *rq, struct task_struct *p, int queued);
34     void (*task_fork)(struct task_struct *p);
35     void (*task_dead)(struct task_struct *p);
36
37     /*
38      * The switched_from() call is allowed to drop rq->lock, therefore we
39      * cannot assume the switched_from/switched_to pair is serialized by
40      * rq->lock. They are however serialized by p->pi_lock.
41      */
42     void (*switched_from)(struct rq *this_rq, struct task_struct *task);
43     void (*switched_to) (struct rq *this_rq, struct task_struct *task);
44     void (*prio_changed) (struct rq *this_rq, struct task_struct *task,
45                          int oldprio);
46
47     unsigned int (*get_rr_interval)(struct rq *rq,
48                                    struct task_struct *task);
49
50     void (*update_curr)(struct rq *rq);
51
52 #define TASK_SET_GROUP      0
53 #define TASK_MOVE_GROUP    1
54
55 #ifdef CONFIG_FAIR_GROUP_SCHED
56     void (*task_change_group)(struct task_struct *p, int type);
57 #endif
58 } __aligned(STRUCT_ALIGNMENT); /* STRUCT_ALIGN(), vmlinux.lds.h */

```

## 3.2 CFS调度策略类

kernel/sched/fair.c 中定义并初始化了完全公平调度策略的调度类 fair\_sched\_class

```

1  /*
2   * All the scheduling class methods:
3   */
4  const struct sched_class fair_sched_class
5  __attribute__((section("__fair_sched_class"))) = {
6      .enqueue_task      = enqueue_task_fair,
7      .dequeue_task      = dequeue_task_fair,
8      .yield_task        = yield_task_fair,
9      .yield_to_task     = yield_to_task_fair,
10
11      .check_preempt_curr = check_preempt_wakeup,
12
13      .pick_next_task     = __pick_next_task_fair,
14      .put_prev_task      = put_prev_task_fair,
15      .set_next_task      = set_next_task_fair,
16

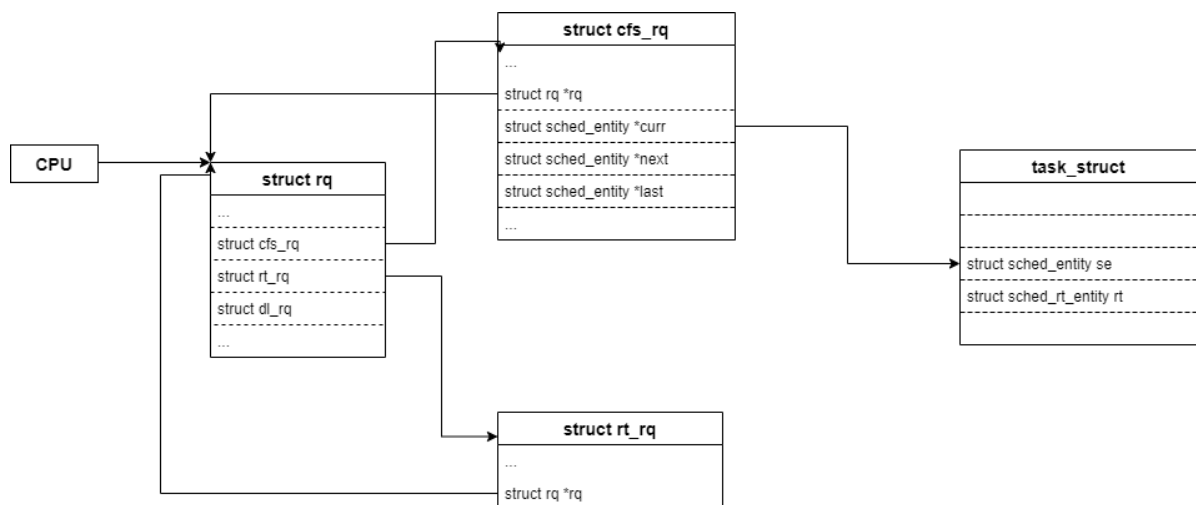
```

```

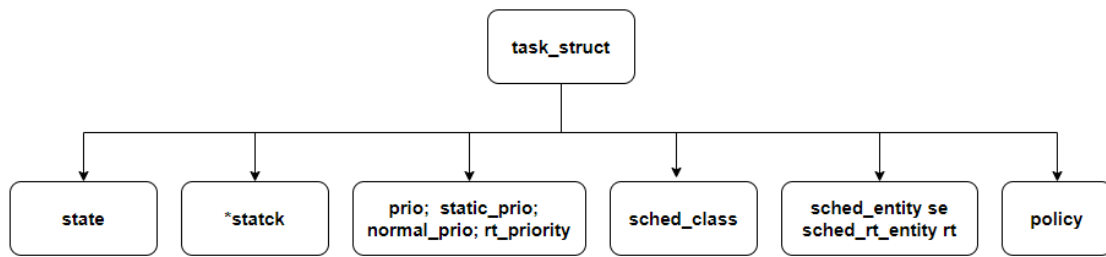
17 #ifdef CONFIG_SMP
18     .balance          = balance_fair,
19     .select_task_rq   = select_task_rq_fair,
20     .migrate_task_rq  = migrate_task_rq_fair,
21
22     .rq_online        = rq_online_fair,
23     .rq_offline       = rq_offline_fair,
24
25     .task_dead        = task_dead_fair,
26     .set_cpus_allowed = set_cpus_allowed_common,
27 #endif
28
29     .task_tick        = task_tick_fair,
30     .task_fork        = task_fork_fair,
31
32     .prio_changed     = prio_changed_fair,
33     .switched_from    = switched_from_fair,
34     .switched_to      = switched_to_fair,
35
36     .get_rr_interval  = get_rr_interval_fair,
37
38     .update_curr      = update_curr_fair,
39
40 #ifdef CONFIG_FAIR_GROUP_SCHED
41     .task_change_group = task_change_group_fair,
42 #endif
43
44 #ifdef CONFIG_UCLAMP_TASK
45     .uclamp_enabled    = 1,
46 #endif
47 };

```

## 四、进程描述符



task\_struct中包含了很多重要的元素，如进程状态、栈内存指针、进程优先级（动态、静态、普通、实时）、进程所属调度类、进程调度实体（普通和实时）、调度策略等等。



`task_struct` 在 `include/linux/sched.h` 中声明，值得好好分析一下。

```

1  struct task_struct {
2  #ifdef CONFIG_THREAD_INFO_IN_TASK
3      /*
4       * For reasons of header soup (see current_thread_info()), this
5       * must be the first element of task_struct.
6       */
7      struct thread_info      thread_info;
8  #endif
9      /* -1 unrunnable, 0 runnable, >0 stopped: */
10     volatile long            state;                                /*
11                                                                    -1表示不可运行，0表示可运行，>0表示停止 */
12     /*
13      * This begins the randomizable portion of task_struct. Only
14      * scheduling-critical items should be added above here.
15      */
16     randomized_struct_fields_start
17
18     void                        *stack;
19     refcount_t                  usage;
20     /* Per task flags (PF_*), defined further below: */
21     unsigned int                flags;                                /*
22                                                                    每进程标志，下文定义 */
23     unsigned int                ptrace;
24
25 #ifdef CONFIG_SMP
26     int                        on_cpu;
27     struct __call_single_node    wake_entry;
28 #ifdef CONFIG_THREAD_INFO_IN_TASK
29     /* Current CPU: */
30     unsigned int                cpu;
31 #endif
32     unsigned int                wakee_flips;
33     unsigned long                wakee_flip_decay_ts;
34     struct task_struct            *last_wakee;
35
36     /*
37      * recent_used_cpu is initially set as the last CPU used by a task
38      * that wakes affine another task. waker/wakee relationships can
39      * push tasks around a CPU where each wakeup moves to the next one.
40      * Tracking a recently used CPU allows a quick search for a recently
41      * used CPU that may be idle.
42      */

```

```

42     int                recent_used_cpu;
43     int                wake_cpu;
44 #endif
45     int                on_rq;
46
47     int                prio;                                /*
动态优先级 */
48     int                static_prio;                        /*
静态优先级 */
49     int                normal_prio;                        /*
普通优先级 */
50     unsigned int       rt_priority;                        /*
实时优先级 */
51
52     const struct sched_class *sched_class;                /*
每个PCB都会指向一个调度类，表示该进程所属的调度器类*/
53     struct sched_entity se;                                /*
普通调度实体 */
54     struct sched_rt_entity rt;                            /*
实时调度实体 */
55 #ifdef CONFIG_CGROUP_SCHED
56     struct task_group   *sched_task_group;
57 #endif
58     struct sched_dl_entity dl;                            /*
空闲调度实体 */
59
60 #ifdef CONFIG_UCLAMP_TASK
61     /*
62      * Clamp values requested for a scheduling entity.
63      * Must be updated with task_rq_lock() held.
64      */
65     struct uclamp_se     uclamp_req[UCLAMP_CNT];
66     /*
67      * Effective clamp values used for a scheduling entity.
68      * Must be updated with task_rq_lock() held.
69      */
70     struct uclamp_se     uclamp[UCLAMP_CNT];
71 #endif
72
73 #ifdef CONFIG_PREEMPT_NOTIFIERS
74     /* List of struct preempt_notifier: */
75     struct hlist_head    preempt_notifiers;
76 #endif
77
78 #ifdef CONFIG_BLK_DEV_IO_TRACE
79     unsigned int         btrace_seq;
80 #endif
81
82     unsigned int         policy;
83     int                  nr_cpus_allowed;
84     const cpumask_t      *cpus_ptr;
85     cpumask_t            cpus_mask;
86
87 #ifdef CONFIG_PREEMPT_RCU
88     int                  rcu_read_lock_nesting;
89     union rcu_special     rcu_read_unlock_special;
90     struct list_head      rcu_node_entry;
91     struct rcu_node        *rcu_blocked_node;

```

```

92 #endif /* #ifdef CONFIG_PREEMPT_RCU */
93
94 #ifdef CONFIG_TASKS_RCU
95     unsigned long        rcu_tasks_nvcsw;
96     u8                   rcu_tasks_holdout;
97     u8                   rcu_tasks_idx;
98     int                  rcu_tasks_idle_cpu;
99     struct list_head      rcu_tasks_holdout_list;
100 #endif /* #ifdef CONFIG_TASKS_RCU */
101
102 #ifdef CONFIG_TASKS_TRACE_RCU
103     int                  trc_reader_nesting;
104     int                  trc_ipi_to_cpu;
105     union rcu_special      trc_reader_special;
106     bool                 trc_reader_checked;
107     struct list_head      trc_holdout_list;
108 #endif /* #ifdef CONFIG_TASKS_TRACE_RCU */
109
110     struct sched_info      sched_info;
111
112     struct list_head      tasks;
113 #ifdef CONFIG_SMP
114     struct plist_node      pushable_tasks;
115     struct rb_node         pushable_dl_tasks;
116 #endif
117
118     struct mm_struct       *mm;
119     struct mm_struct       *active_mm;
120
121     /* Per-thread vma caching: */
122     struct vmacache        vmacache;
123
124 #ifdef SPLIT_RSS_COUNTING
125     struct task_rss_stat    rss_stat;
126 #endif
127     int                    exit_state;
128     int                    exit_code;
129     int                    exit_signal;
130     /* The signal sent when the parent dies: */
131     int                    pdeath_signal;
132     /* JOBCTL_*, siglock protected: */
133     unsigned long          jobctl;
134
135     /* Used for emulating ABI behavior of previous Linux versions: */
136     unsigned int           personality;
137
138     /* Scheduler bits, serialized by scheduler locks: */
139     unsigned               sched_reset_on_fork:1;
140     unsigned               sched_contributes_to_load:1;
141     unsigned               sched_migrated:1;
142     unsigned               sched_remote_wakeup:1;
143 #ifdef CONFIG_PSI
144     unsigned               sched_psi_wake_requeue:1;
145 #endif
146
147     /* Force alignment to the next boundary: */
148     unsigned               :0;
149

```

```

150     /* Unserialized, strictly 'current' */
151
152     /* Bit to tell LSMs we're in execve(): */
153     unsigned         in_execve:1;
154     unsigned         in_iowait:1;
155 #ifndef TIF_RESTORE_SIGMASK
156     unsigned         restore_sigmask:1;
157 #endif
158 #ifdef CONFIG_MEMCG
159     unsigned         in_user_fault:1;
160 #endif
161 #ifdef CONFIG_COMPAT_BRK
162     unsigned         brk_randomized:1;
163 #endif
164 #ifdef CONFIG_CGROUPS
165     /* disallow userland-initiated cgroup migration */
166     unsigned         no_cgroup_migration:1;
167     /* task is frozen/stopped (used by the cgroup freezer) */
168     unsigned         frozen:1;
169 #endif
170 #ifdef CONFIG_BLK_CGROUP
171     unsigned         use_memdelay:1;
172 #endif
173 #ifdef CONFIG_PSI
174     /* Stalled due to lack of memory */
175     unsigned         in_memstall:1;
176 #endif
177
178     unsigned long     atomic_flags; /* Flags requiring atomic access.
179     */
180
181     struct restart_block restart_block;
182
183     pid_t             pid;
184     pid_t             tgid;
185
186 #ifdef CONFIG_STACKPROTECTOR
187     /* Canary value for the -fstack-protector GCC feature: */
188     unsigned long     stack_canary;
189 #endif
190     /*
191      * Pointers to the (original) parent process, youngest child, younger
192      * sibling,
193      * older sibling, respectively. (p->father can be replaced with
194      * p->real_parent->pid)
195      */
196     /* Real parent process: */
197     struct task_struct __rcu *real_parent;
198
199     /* Recipient of SIGCHLD, wait4() reports: */
200     struct task_struct __rcu *parent;
201
202     /*
203      * Children/sibling form the list of natural children:
204      */
205     struct list_head   children;
206     struct list_head   sibling;

```

```

206     struct task_struct      *group_leader;
207
208     /*
209      * 'ptraced' is the list of tasks this task is using ptrace() on.
210      *
211      * This includes both natural children and PTRACE_ATTACH targets.
212      * 'ptrace_entry' is this task's link on the p->parent->ptraced list.
213      */
214     struct list_head        ptraced;
215     struct list_head        ptrace_entry;
216
217     /* PID/PID hash table linkage. */
218     struct pid              *thread_pid;
219     struct hlist_node        pid_links[PIDTYPE_MAX];
220     struct list_head        thread_group;
221     struct list_head        thread_node;
222
223     struct completion        *vfork_done;
224
225     /* CLONE_CHILD_SETTID: */
226     int __user               *set_child_tid;
227
228     /* CLONE_CHILD_CLEARTID: */
229     int __user               *clear_child_tid;
230
231     u64                      utime;
232     u64                      stime;
233 #ifdef CONFIG_ARCH_HAS_SCALED_CPUTIME
234     u64                      utimescaled;
235     u64                      stimescaled;
236 #endif
237     u64                      gtime;
238     struct prev_cputime      prev_cputime;
239 #ifdef CONFIG_VIRT_CPU_ACCOUNTING_GEN
240     struct vtime             vtime;
241 #endif
242
243 #ifdef CONFIG_NO_HZ_FULL
244     atomic_t                 tick_dep_mask;
245 #endif
246     /* Context switch counts: */
247     unsigned long            nvcs;
248     unsigned long            nivcs;
249
250     /* Monotonic time in nsecs: */
251     u64                      start_time;
252
253     /* Boot based time in nsecs: */
254     u64                      start_boottime;
255
256     /* MM fault and swap info: this can arguably be seen as either mm-
257      specific or thread-specific: */
257     unsigned long            min_flt;
258     unsigned long            maj_flt;
259
260     /* Empty if CONFIG_POSIX_CPUTIMERS=n */
261     struct posix_cputimers    posix_cputimers;
262

```



```

263 #ifdef CONFIG_POSIX_CPU_TIMERS_TASK_WORK
264     struct posix_cputimers_work posix_cputimers_work;
265 #endif
266
267     /* Process credentials: */
268
269     /* Tracer's credentials at attach: */
270     const struct cred __rcu    *ptracer_cred;
271
272     /* Objective and real subjective task credentials (COW): */
273     const struct cred __rcu    *real_cred;
274
275     /* Effective (overridable) subjective task credentials (COW): */
276     const struct cred __rcu    *cred;
277
278 #ifdef CONFIG_KEYS
279     /* Cached requested key. */
280     struct key                 *cached_requested_key;
281 #endif
282
283     /*
284      * executable name, excluding path.
285      *
286      * - normally initialized setup_new_exec()
287      * - access it with [gs]et_task_comm()
288      * - lock it with task_lock()
289      */
290     char                       comm[TASK_COMM_LEN];
291
292     struct nameidata           *nameidata;
293
294 #ifdef CONFIG_SYSVIPC
295     struct sysv_sem            sysvsem;
296     struct sysv_shm            sysvshm;
297 #endif
298 #ifdef CONFIG_DETECT_HUNG_TASK
299     unsigned long              last_switch_count;
300     unsigned long              last_switch_time;
301 #endif
302     /* Filesystem information: */
303     struct fs_struct           *fs;
304
305     /* Open file information: */
306     struct files_struct        *files;
307
308 #ifdef CONFIG_IO_URING
309     struct io_uring_task       *io_uring;
310 #endif
311
312     /* Namespaces: */
313     struct nsproxy             *nsproxy;
314
315     /* Signal handlers: */
316     struct signal_struct       *signal;
317     struct sighand_struct __rcu *sighand;
318     sigset_t                   blocked;
319     sigset_t                   real_blocked;
320     /* Restored if set_restore_sigmask() was used: */

```

```

321     sigset_t          saved_sigmask;
322     struct sigpending    pending;
323     unsigned long      sas_ss_sp;
324     size_t             sas_ss_size;
325     unsigned int        sas_ss_flags;
326
327     struct callback_head *task_works;
328
329 #ifdef CONFIG_AUDIT
330 #ifdef CONFIG_AUDITSYSCALL
331     struct audit_context *audit_context;
332 #endif
333     kuid_t              loginuid;
334     unsigned int         sessionid;
335 #endif
336     struct seccomp        seccomp;
337
338     /* Thread group tracking: */
339     u64                  parent_exec_id;
340     u64                  self_exec_id;
341
342     /* Protection against (de-)allocation: mm, files, fs, tty, keyrings,
343     mems_allowed, mempolicy: */
344     spinlock_t           alloc_lock;
345
346     /* Protection of the PI data structures: */
347     raw_spinlock_t       pi_lock;
348
349     struct wake_q_node    wake_q;
350
351 #ifdef CONFIG_RT_MUTEXES
352     /* PI waiters blocked on a rt_mutex held by this task: */
353     struct rb_root_cached pi_waiters;
354     /* Updated under owner's pi_lock and rq lock */
355     struct task_struct    *pi_top_task;
356     /* Deadlock detection and priority inheritance handling: */
357     struct rt_mutex_waiter *pi_blocked_on;
358 #endif
359
360 #ifdef CONFIG_DEBUG_MUTEXES
361     /* Mutex deadlock detection: */
362     struct mutex_waiter    *blocked_on;
363 #endif
364
365 #ifdef CONFIG_DEBUG_ATOMIC_SLEEP
366     int                    non_block_count;
367 #endif
368
369 #ifdef CONFIG_TRACE_IRQFLAGS
370     struct irqtrace_events irqtrace;
371     unsigned int           hardirq_threaded;
372     u64                    hardirq_chain_key;
373     int                    softirqs_enabled;
374     int                    softirq_context;
375     int                    irq_config;
376 #endif
377 #ifdef CONFIG_LOCKDEP

```

```

378 # define MAX_LOCK_DEPTH          48UL
379     u64                curr_chain_key;
380     int                lockdep_depth;
381     unsigned int       lockdep_recursion;
382     struct held_lock    held_locks[MAX_LOCK_DEPTH];
383 #endif
384
385 #ifdef CONFIG_UBSAN
386     unsigned int       in_ubsan;
387 #endif
388
389     /* Journalling filesystem info: */
390     void                *journal_info;
391
392     /* Stacked block device info: */
393     struct bio_list     *bio_list;
394
395 #ifdef CONFIG_BLOCK
396     /* Stack plugging: */
397     struct blk_plug     *plug;
398 #endif
399
400     /* VM state: */
401     struct reclaim_state *reclaim_state;
402
403     struct backing_dev_info *backing_dev_info;
404
405     struct io_context   *io_context;
406
407 #ifdef CONFIG_COMPACTION
408     struct capture_control *capture_control;
409 #endif
410     /* Ptrace state: */
411     unsigned long        ptrace_message;
412     kernel_siginfo_t     *last_siginfo;
413
414     struct task_io_accounting ioac;
415 #ifdef CONFIG_PSI
416     /* Pressure stall state */
417     unsigned int         psi_flags;
418 #endif
419 #ifdef CONFIG_TASK_XACCT
420     /* Accumulated RSS usage: */
421     u64                  acct_rss_mem1;
422     /* Accumulated virtual memory usage: */
423     u64                  acct_vm_mem1;
424     /* stime + utime since last update: */
425     u64                  acct_timexpd;
426 #endif
427 #ifdef CONFIG_CPUSETS
428     /* Protected by ->alloc_lock: */
429     nodemask_t           mems_allowed;
430     /* Sequence number to catch updates: */
431     seqcount_spinlock_t  mems_allowed_seq;
432     int                  cpuset_mem_spread_rotor;
433     int                  cpuset_slab_spread_rotor;
434 #endif
435 #ifdef CONFIG_CGROUPS

```

```

436  /* Control Group info protected by css_set_lock: */
437  struct css_set __rcu      *cgroups;
438  /* cg_list protected by css_set_lock and tsk->alloc_lock: */
439  struct list_head         cg_list;
440 #endif
441 #ifdef CONFIG_X86_CPU_RESCTRL
442     u32                     closid;
443     u32                     rmid;
444 #endif
445 #ifdef CONFIG_FUTEX
446     struct robust_list_head __user *robust_list;
447 #ifdef CONFIG_COMPAT
448     struct compat_robust_list_head __user *compat_robust_list;
449 #endif
450     struct list_head        pi_state_list;
451     struct futex_pi_state   *pi_state_cache;
452     struct mutex            futex_exit_mutex;
453     unsigned int            futex_state;
454 #endif
455 #ifdef CONFIG_PERF_EVENTS
456     struct perf_event_context *perf_event_ctxp[perf_nr_task_contexts];
457     struct mutex            perf_event_mutex;
458     struct list_head        perf_event_list;
459 #endif
460 #ifdef CONFIG_DEBUG_PREEMPT
461     unsigned long           preempt_disable_ip;
462 #endif
463 #ifdef CONFIG_NUMA
464     /* Protected by alloc_lock: */
465     struct mempolicy        *mempolicy;
466     short                   il_prev;
467     short                   pref_node_fork;
468 #endif
469 #ifdef CONFIG_NUMA_BALANCING
470     int                     numa_scan_seq;
471     unsigned int            numa_scan_period;
472     unsigned int            numa_scan_period_max;
473     int                     numa_preferred_nid;
474     unsigned long           numa_migrate_retry;
475     /* Migration stamp: */
476     u64                     node_stamp;
477     u64                     last_task_numa_placement;
478     u64                     last_sum_exec_runtime;
479     struct callback_head    numa_work;
480
481     /*
482     * This pointer is only modified for current in syscall and
483     * pagefault context (and for tasks being destroyed), so it can be read
484     * from any of the following contexts:
485     * - RCU read-side critical section
486     * - current->numa_group from everywhere
487     * - task's runqueue locked, task not running
488     */
489     struct numa_group __rcu  *numa_group;
490
491     /*
492     * numa_faults is an array split into four regions:
493     * faults_memory, faults_cpu, faults_memory_buffer, faults_cpu_buffer

```

```

494     * in this precise order.
495     *
496     * faults_memory: Exponential decaying average of faults on a per-node
497     * basis. Scheduling placement decisions are made based on these
498     * counts. The values remain static for the duration of a PTE scan.
499     * faults_cpu: Track the nodes the process was running on when a NUMA
500     * hinting fault was incurred.
501     * faults_memory_buffer and faults_cpu_buffer: Record faults per node
502     * during the current scan window. When the scan completes, the counts
503     * in faults_memory and faults_cpu decay and these values are copied.
504     */
505     unsigned long          *numa_faults;
506     unsigned long          total_numa_faults;
507
508     /*
509     * numa_faults_locality tracks if faults recorded during the last
510     * scan window were remote/local or failed to migrate. The task scan
511     * period is adapted based on the locality of the faults with different
512     * weights depending on whether they were shared or private faults
513     */
514     unsigned long          numa_faults_locality[3];
515
516     unsigned long          numa_pages_migrated;
517 #endif /* CONFIG_NUMA_BALANCING */
518
519 #ifdef CONFIG_RSEQ
520     struct rseq __user *rseq;
521     u32 rseq_sig;
522     /*
523     * RMW on rseq_event_mask must be performed atomically
524     * with respect to preemption.
525     */
526     unsigned long rseq_event_mask;
527 #endif
528
529     struct tlbflush_unmap_batch tlb_ubc;
530
531     union {
532         refcount_t      rcu_users;
533         struct rcu_head   rcu;
534     };
535
536     /* Cache last used pipe for splice(): */
537     struct pipe_inode_info *splice_pipe;
538
539     struct page_frag      task_frag;
540
541 #ifdef CONFIG_TASK_DELAY_ACCT
542     struct task_delay_info *delays;
543 #endif
544
545 #ifdef CONFIG_FAULT_INJECTION
546     int          make_it_fail;
547     unsigned int fail_nth;
548 #endif
549     /*
550     * When (nr_dirtied >= nr_dirtied_pause), it's time to call
551     * balance_dirty_pages() for a dirty throttling pause:

```

```

552     */
553     int             nr_dirtied;
554     int             nr_dirtied_pause;
555     /* Start of a write-and-pause period: */
556     unsigned long    dirty_paused_when;
557
558 #ifdef CONFIG_LATENCYTOP
559     int             latency_record_count;
560     struct latency_record latency_record[LT_SAVECOUNT];
561 #endif
562     /*
563      * Time slack values; these are used to round up poll() and
564      * select() etc timeout values. These are in nanoseconds.
565      */
566     u64             timer_slack_ns;
567     u64             default_timer_slack_ns;
568
569 #ifdef CONFIG_KASAN
570     unsigned int     kasan_depth;
571 #endif
572
573 #ifdef CONFIG_KCSAN
574     struct kcsan_ctx kcsan_ctx;
575 #ifdef CONFIG_TRACE_IRQFLAGS
576     struct irqtrace_events kcsan_save_irqtrace;
577 #endif
578 #endif
579
580 #ifdef CONFIG_FUNCTION_GRAPH_TRACER
581     /* Index of current stored address in ret_stack: */
582     int             curr_ret_stack;
583     int             curr_ret_depth;
584
585     /* Stack of return addresses for return function tracing: */
586     struct ftrace_ret_stack *ret_stack;
587
588     /* Timestamp for last schedule: */
589     unsigned long long ftrace_timestamp;
590
591     /*
592      * Number of functions that haven't been traced
593      * because of depth overrun:
594      */
595     atomic_t         trace_overrun;
596
597     /* Pause tracing: */
598     atomic_t         tracing_graph_pause;
599 #endif
600
601 #ifdef CONFIG_TRACING
602     /* State flags for use by tracers: */
603     unsigned long    trace;
604
605     /* Bitmask and counter of trace recursion: */
606     unsigned long    trace_recursion;
607 #endif /* CONFIG_TRACING */
608
609 #ifdef CONFIG_KCOV

```

```

610     /* See kernel/kcov.c for more details. */
611
612     /* Coverage collection mode enabled for this task (0 if disabled): */
613     unsigned int          kcov_mode;
614
615     /* Size of the kcov_area: */
616     unsigned int          kcov_size;
617
618     /* Buffer for coverage collection: */
619     void                  *kcov_area;
620
621     /* KCOV descriptor wired with this task or NULL: */
622     struct kcov            *kcov;
623
624     /* KCOV common handle for remote coverage collection: */
625     u64                    kcov_handle;
626
627     /* KCOV sequence number: */
628     int                    kcov_sequence;
629
630     /* Collect coverage from softirq context: */
631     unsigned int          kcov_softirq;
632 #endif
633
634 #ifdef CONFIG_MEMCG
635     struct mem_cgroup      *memcg_in_oom;
636     gfp_t                  memcg_oom_gfp_mask;
637     int                    memcg_oom_order;
638
639     /* Number of pages to reclaim on returning to userland: */
640     unsigned int          memcg_nr_pages_over_high;
641
642     /* Used by memcontrol for targeted memcg charge: */
643     struct mem_cgroup      *active_memcg;
644 #endif
645
646 #ifdef CONFIG_BLK_CGROUP
647     struct request_queue   *throttle_queue;
648 #endif
649
650 #ifdef CONFIG_UPROBES
651     struct uprobe_task     *utask;
652 #endif
653 #if defined(CONFIG_BCACHE) || defined(CONFIG_BCACHE_MODULE)
654     unsigned int          sequential_io;
655     unsigned int          sequential_io_avg;
656 #endif
657 #ifdef CONFIG_DEBUG_ATOMIC_SLEEP
658     unsigned long          task_state_change;
659 #endif
660     int                    pagefault_disabled;
661 #ifdef CONFIG_MMU
662     struct task_struct      *oom_reaper_list;
663 #endif
664 #ifdef CONFIG_VMAP_STACK
665     struct vm_struct        *stack_vm_area;
666 #endif
667 #ifdef CONFIG_THREAD_INFO_IN_TASK

```

```

668     /* A live task holds one reference: */
669     refcount_t          stack_refcount;
670 #endif
671 #ifdef CONFIG_LIVEPATCH
672     int patch_state;
673 #endif
674 #ifdef CONFIG_SECURITY
675     /* Used by LSM modules for access restriction: */
676     void                *security;
677 #endif
678
679 #ifdef CONFIG_GCC_PLUGIN_STACKLEAK
680     unsigned long        lowest_stack;
681     unsigned long        prev_lowest_stack;
682 #endif
683
684 #ifdef CONFIG_X86_MCE
685     u64                  mce_addr;
686     __u64                 mce_ripv : 1,
687                        mce_whole_page : 1,
688                        __mce_reserved : 62;
689     struct callback_head  mce_kill_me;
690 #endif
691
692     /*
693      * New fields for task_struct should be added above here, so that
694      * they are included in the randomized portion of task_struct.
695      */
696     randomized_struct_fields_end
697
698     /* CPU-specific state of this task: */
699     struct thread_struct   thread;
700
701     /*
702      * WARNING: on x86, 'thread_struct' contains a variable-sized
703      * structure. It *MUST* be at the end of 'task_struct'.
704      *
705      * Do not put anything below here!
706      */
707 }

```

## 五、调度器

### 5.1 周期性调度器

周期性调度器在scheduler\_tick中实现。如果系统正在活动中，内核会按照频率HZ自动调用该函数。如果没有进程在等待调度，那么在计算机电力供应不足的情况下，也可以关闭该调度器以减少电能消耗。

kernel/sched/core.c 中定义

```

1  /*
2   * This function gets called by the timer code, with HZ frequency.
3   * We call it with interrupts disabled.
4   */
5  void scheduler_tick(void)

```



```

6  {
7      int cpu = smp_processor_id();                /* 获取当前
cpu号 */
8      struct rq *rq = cpu_rq(cpu);                /* 获取cpu就
绪队列rq（每个cpu都有一个就绪队列） */
9      struct task_struct *curr = rq->curr;        /* 从rq中获取
当前运行进程的描述符 */
10     struct rq_flags rf;
11     unsigned long thermal_pressure;              /* 5.7内核后
的新特性，CPU热压过高后会限频，但调度器并不知道，所以需要让调度器感知CPU频率被限制住，这样
更好的调度任务*/
12
13     arch_scale_freq_tick();
14     sched_clock_tick();
15
16     rq_lock(rq, &rf);
17
18     update_rq_clock(rq);                          /* 更新就绪队列中的clock和clock_task成员
值，代表当前的时间，一般我们会用到clock_task*/
19     thermal_pressure = arch_scale_thermal_pressure(cpu_of(rq));
20     update_thermal_load_avg(rq_clock_thermal(rq), rq, thermal_pressure);
21     curr->sched_class->task_tick(rq, curr, 0);    /* 进入当前进程的调度类的
task_tick函数中，更新当前进程的时间片，不同调度类的该函数实现不同*/
22     calc_global_load_tick(rq);
23     psi_task_tick(rq);
24
25     rq_unlock(rq, &rf);
26
27     perf_event_task_tick();
28
29 #ifdef CONFIG_SMP
30     rq->idle_balance = idle_cpu(cpu);             /* 判断cpu是否空闲 */
31     trigger_load_balance(rq);                    /* 挂起SCHED_SOFTIRQ软中断函
数，去做周期性的负载平衡操作 */
32 #endif
33 }

```

该函数主要作用：

- (1) 管理内核中与整个系统和各个进程的调度相关的统计量。其间执行的主要操作是对各种计数器加1，我们对此没什么兴趣。
- (2) 激活负责当前进程的调度类的周期性调度方法。

## 5.2 主调度器

在内核中的许多地方，如果要将CPU分配给与当前活动进程不同的另一个进程，都会直接调用主调度器函数（schedule）。在从系统调用返回之后，内核也会检查当前进程是否设置了重调度标志TIF\_NEED\_RESCHED，例如，前述的scheduler\_tick就会设置该标志。如果是这样，则内核会调用schedule。该函数假定当前活动进程一定会被另一个进程取代。

### 5.2.1 \_\_sched前缀

```
1  /* Attach to any functions which should be ignored in wchan output. */
2  #define __sched      __attribute__((__section__(".sched.text")))
```

将相关函数的代码编译之后，放到目标文件的一个特定的段中，即 `.sched.text` 中。该信息使得内核在显示栈转储或类似信息时，忽略所有与调度有关的调用。

### 5.2.2 schedule函数

```
1  #define tif_need_resched() test_thread_flag(TIF_NEED_RESCHED)
2
3  static __always_inline bool need_resched(void)      /* 该函数用于判断
   TIF_NEED_RESCHED标志位看是否需要重新调度 */
4  {
5      return unlikely(tif_need_resched());
6  }
```

#### schedule() 函数

```
1  asmlinkage __visible void __sched schedule(void)
2  {
3      struct task_struct *tsk = current;           /* 获取当前进程的结构体*/
4
5      sched_submit_work(tsk);                       /* 防止死锁问题 */
6      do {
7          preempt_disable();                       /* 关闭抢占 */
8          __schedule(false);
9          sched_preempt_enable_no_resched();        /* 开启抢占 */
10     } while (need_resched());                     /* 如果需要重新调度，则循环? */
11     sched_update_worker(tsk);
12 }
13 EXPORT_SYMBOL(schedule);
```

#### \_\_schedule() 函数

```
1  /*
2   * __schedule()是主要的调度函数。
3   * 所谓的主要函数是指推动调度，因此进入该函数的原因有：
4   * 1. 明显的阻塞：锁、信号、等待队列等等；
5   * 2. TIF_NEED_RESCHED标志被中断和用户空间返回路径检测到；
6   * 3. 唤醒wakeup并没有真正的进入schedule()，唤醒只是将进程加入run-queue
7   * The main means of driving the scheduler and thus entering this function
   are:
8   *
9   * 1. Explicit blocking: mutex, semaphore, waitqueue, etc.
10  *
11  * 2. TIF_NEED_RESCHED flag is checked on interrupt and userspace return
12  *    paths. For example, see arch/x86/entry_64.S.
13  *
14  * To drive preemption between tasks, the scheduler sets the flag in
   timer
15  * interrupt handler scheduler_tick().
16  */
```

```

17  * 3. wakeups don't really cause entry into schedule(). They add a
18  * task to the run-queue and that's it.
19  *
20  * Now, if the new task added to the run-queue preempts the current
21  * task, then the wakeup sets TIF_NEED_RESCHED and schedule() gets
22  * called on the nearest possible occasion:
23  *
24  * - If the kernel is preemptible (CONFIG_PREEMPTION=y):
25  *
26  * - in syscall or exception context, at the next outmost
27  * preempt_enable(). (this might be as soon as the wake_up()'s
28  * spin_unlock()!)
29  *
30  * - in IRQ context, return from interrupt-handler to
31  * preemptible context
32  *
33  * - If the kernel is not preemptible (CONFIG_PREEMPTION is not set)
34  * then at the next:
35  *
36  * - cond_resched() call
37  * - explicit schedule() call
38  * - return from syscall or exception to user-space
39  * - return from interrupt-handler to user-space
40  *
41  * WARNING: must be called with preemption disabled! 调用时必须禁用抢占
42  */
43 static void __sched notrace __schedule(bool preempt)
44 {
45     struct task_struct *prev, *next;
46     unsigned long *switch_count;
47     unsigned long prev_state;
48     struct rq_flags rf;
49     struct rq *rq;
50     int cpu;
51
52     cpu = smp_processor_id(); /* 获取当前cpu号 */
53     rq = cpu_rq(cpu); /* 获取当前cpu的
runqueue */
54     prev = rq->curr; /* 将当前进程的描述符
指针保存在prev变量中 */
55
56     schedule_debug(prev, preempt);
57
58     if (sched_feat(HRTICK))
59         hrtick_clear(rq);
60
61     local_irq_disable();
62     rcu_note_context_switch(preempt);
63
64     /*
65      * Make sure that signal_pending_state()->signal_pending() below
66      * can't be reordered with __set_current_state(TASK_INTERRUPTIBLE)
67      * done by the caller to avoid the race with signal_wake_up():
68      *
69      * __set_current_state(@state)      signal_wake_up()
70      * schedule()                      set_tsk_thread_flag(p, TIF_SIGPENDING)
71      *                                wake_up_state(p, state)
72      * LOCK rq->lock                    LOCK p->pi_state

```

```

73      *   smp_mb__after_spinlock()          smp_mb__after_spinlock()
74      *   if (signal_pending_state())      if (p->state & @state)
75      *
76      * Also, the membarrier system call requires a full memory barrier
77      * after coming from user-space, before storing to rq->curr.
78      */
79  rq_lock(rq, &rf);
80  smp_mb__after_spinlock();
81
82  /* Promote REQ to ACT */
83  rq->clock_update_flags <=< 1;
84  update_rq_clock(rq);
85
86  switch_count = &prev->nivcsw;
87
88  /*
89   * We must load prev->state once (task_struct::state is volatile), such
90   * that:
91   *
92   * - we form a control dependency vs deactivate_task() below.
93   * - ptrace_{,un}freeze_traced() can change ->state underneath us.
94   */
95  prev_state = prev->state;
96  if (!preempt && prev_state) {
97      if (signal_pending_state(prev_state, prev)) {
98          prev->state = TASK_RUNNING;
99      } else {
100          prev->sched_contributes_to_load =
101              (prev_state & TASK_UNINTERRUPTIBLE) &&
102              !(prev_state & TASK_NOLOAD) &&
103              !(prev->flags & PF_FROZEN);
104
105          if (prev->sched_contributes_to_load)
106              rq->nr_uninterruptible++;
107
108          /*
109           * __schedule()          ttwu()
110           *   prev_state = prev->state;   if (p->on_rq && ...)
111           *   if (prev_state)             goto out;
112           *   p->on_rq = 0;                smp_acquire__after_ctrl_dep();
113           *                               p->state = TASK_WAKING
114           *
115           * where __schedule() and ttwu() have matching control
116           dependencies.
117           *
118           * After this, schedule() must not care about p->state any
119           more.
120           */
121          deactivate_task(rq, prev, DEQUEUE_SLEEP | DEQUEUE_NOCLOCK);
122
123          if (prev->in_iowait) {
124              atomic_inc(&rq->nr_iowait);
125              delayacct_blkio_start();
126          }
127      }
128      switch_count = &prev->nvcs;
129  }

```

```

129     next = pick_next_task(rq, prev, &rf);          /* 将下一个被调度的进程描述
符指针存放在next变量中 */
130     clear_tsk_need_resched(prev);                  /* 清除当前进程的
TIF_NEED_RESCHED标志位 */
131     clear_preempt_need_resched();                  /* 清除
PREEMPT_NEED_RESCHED */
132
133     if (likely(prev != next)) {
134         rq->nr_switches++;
135         /*
136          * RCU users of rcu_dereference(rq->curr) may not see
137          * changes to task_struct made by pick_next_task().
138          */
139         RCU_INIT_POINTER(rq->curr, next);
140         /*
141          * The membarrier system call requires each architecture
142          * to have a full memory barrier after updating
143          * rq->curr, before returning to user-space.
144          *
145          * Here are the schemes providing that barrier on the
146          * various architectures:
147          * - mm ? switch_mm() : mmdrop() for x86, s390, sparc, PowerPC.
148          *   switch_mm() rely on membarrier_arch_switch_mm() on PowerPC.
149          * - finish_lock_switch() for weakly-ordered
150          *   architectures where spin_unlock is a full barrier,
151          * - switch_to() for arm64 (weakly-ordered, spin_unlock
152          *   is a RELEASE barrier),
153          */
154         ++*switch_count;
155
156         psi_sched_switch(prev, next, !task_on_rq_queued(prev));
157
158         trace_sched_switch(preempt, prev, next);          /*
event事件 sched_switch */
159
160         /* Also unlocks the rq: */
161         rq = context_switch(rq, prev, next, &rf);          /* 当前进程和下一
个进程的上下文进行切换*/
162     } else {
163         rq->clock_update_flags &= ~(RQCF_ACT_SKIP|RQCF_REQ_SKIP);
164         rq_unlock_irq(rq, &rf);
165     }
166
167     balance_callback(rq);
168 }

```

## 上下文切换context\_switch()

上下文切换一般分为两个，一个是硬件上下文切换（指的是cpu寄存器，要把当前进程使用的寄存器内容保存下来，再把下一个程序的寄存器内容恢复），另一个是切换进程的地址空间（说白了就是程序代码）。进程的地址空间（程序代码）主要保存在进程描述符中struct mm\_struct结构体中，因此该函数主要是操作这个结构体。

```

1  /*
2  * context_switch - switch to the new MM and the new thread's register
state.
3  */

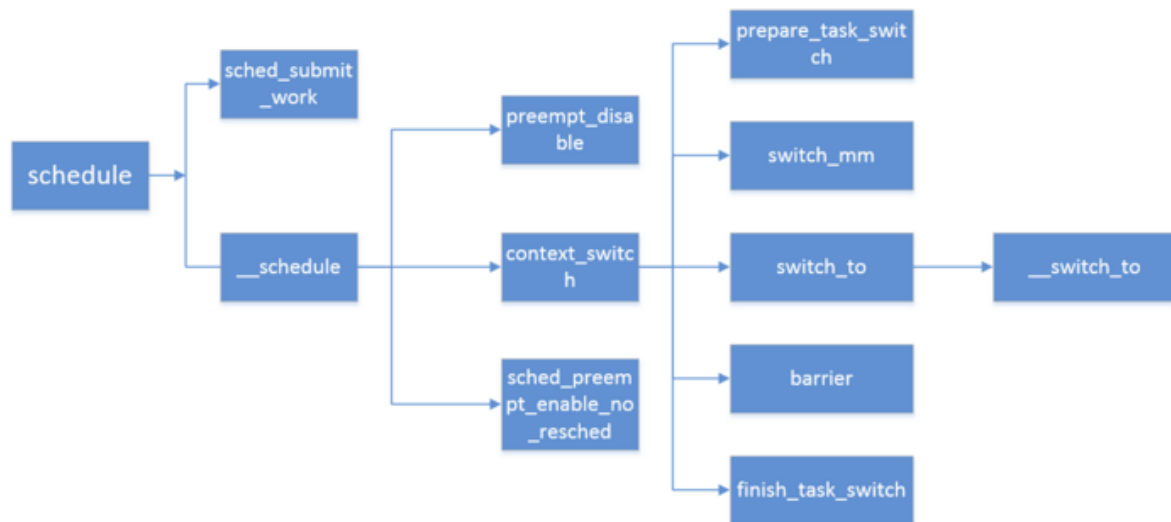
```

```

4 static __always_inline struct rq *
5 context_switch(struct rq *rq, struct task_struct *prev,
6               struct task_struct *next, struct rq_flags *rf)
7 {
8     prepare_task_switch(rq, prev, next);
9
10    /*
11     * For paravirt, this is coupled with an exit in switch_to to
12     * combine the page table reload and the switch backend into
13     * one hypercall.
14     */
15    arch_start_context_switch(prev);
16
17    /*
18     * kernel -> kernel    lazy + transfer active
19     *   user -> kernel    lazy + mmgrab() active
20     *
21     * kernel ->   user    switch + mmdrop() active
22     *   user ->   user    switch
23     */
24    if (!next->mm) {                                     // to kernel
25        enter_lazy_tlb(prev->active_mm, next);
26
27        next->active_mm = prev->active_mm;
28        if (prev->mm)                                     // from user
29            mmgrab(prev->active_mm);
30        else
31            prev->active_mm = NULL;
32    } else {                                             // to user
33        membarrier_switch_mm(rq, prev->active_mm, next->mm);
34        /*
35         * sys_membarrier() requires an smp_mb() between setting
36         * rq->curr / membarrier_switch_mm() and returning to userspace.
37         *
38         * The below provides this either through switch_mm(), or in
39         * case 'prev->active_mm == next->mm' through
40         * finish_task_switch()'s mmdrop().
41         */
42        switch_mm_irqs_off(prev->active_mm, next->mm, next);
43
44        if (!prev->mm) {                                 // from kernel
45            /* will mmdrop() in finish_task_switch(). */
46            rq->prev_mm = prev->active_mm;
47            prev->active_mm = NULL;
48        }
49    }
50
51    rq->clock_update_flags &= ~(RQCF_ACT_SKIP|RQCF_REQ_SKIP);
52
53    prepare_lock_switch(rq, next, rf);
54
55    /* Here we just switch the register state and the stack. */
56    switch_to(prev, next, prev);
57    barrier();
58
59    return finish_task_switch(prev);
60 }

```

网上有个4.x的内核代码，这个和5.9稍微有点不一样，大体流程还是差不多



### 5.2.3 linux-5.9.10调度器框架

调度类的顺序

```
1  /* vmlinux.lds.h文件中定义
2  * The order of the sched class addresses are important, as they are
3  * used to determine the order of the priority of each sched class in
4  * relation to each other.这里固定死了调度类的优先级顺序, stop > deadline > rt >
   fair > idle
5  */
6  #define SCHED_DATA \
7      STRUCT_ALIGN(); \
8      __begin_sched_classes = .; \
9      *(__idle_sched_class) \
10     *(__fair_sched_class) \
11     *(__rt_sched_class) \
12     *(__dl_sched_class) \
13     *(__stop_sched_class) \
14     __end_sched_classes = .;
```

```
1  /*kernel/sched/sched.h中定义*/
2  /* Defined in include/asm-generic/vmlinux.lds.h */
3  extern struct sched_class __begin_sched_classes[];
4  extern struct sched_class __end_sched_classes[];
5
6  #define sched_class_highest (__end_sched_classes - 1)
7  #define sched_class_lowest (__begin_sched_classes - 1) /*此处为何是begin -1
   ? 难道不是 +1 ?*/
8
9  #define for_class_range(class, _from, _to) \
```

```

10     for (class = (_from); class != (_to); class--)
11
12     #define for_each_class(class) \
13         for_class_range(class, sched_class_highest, sched_class_lowest) /*从stop
遍历至idle*/
14
15     extern const struct sched_class stop_sched_class;
16     extern const struct sched_class dl_sched_class;
17     extern const struct sched_class rt_sched_class;
18     extern const struct sched_class fair_sched_class;
19     extern const struct sched_class idle_sched_class;

```

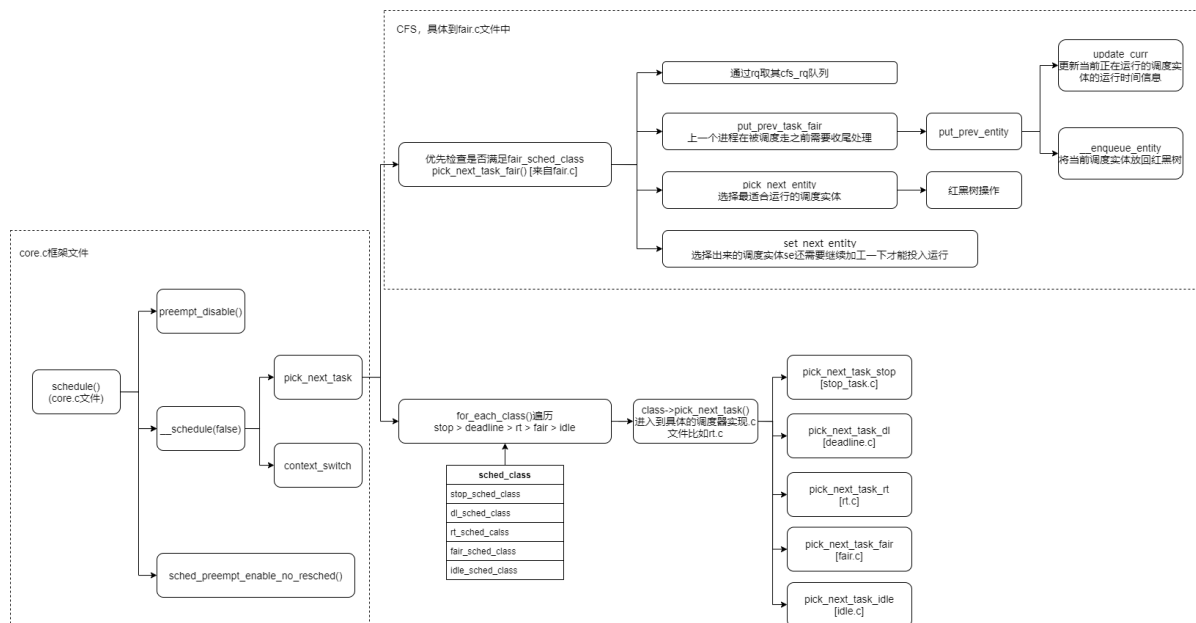


linux5.9.10中调度器的主框架core.c是如何与CFS、实时rt、deadline、idle等具体的调度器实现整合起来的？

```

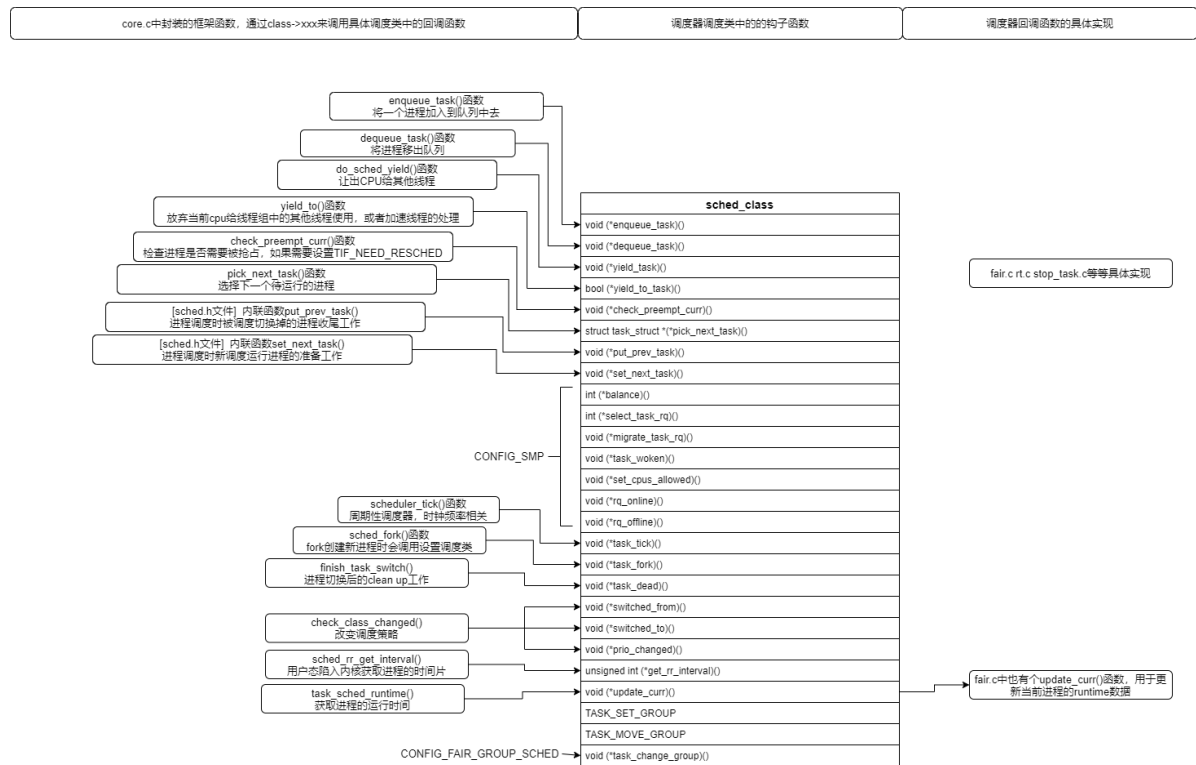
1  /*
2  * Scheduling policies
3  */
4  #define SCHED_NORMAL      0
5  #define SCHED_FIFO        1
6  #define SCHED_RR          2
7  #define SCHED_BATCH       3
8  /* SCHED_ISO: reserved but not implemented yet */ /*MuQSS
使用的是该policy*/
9  #define SCHED_IDLE        5
10 #define SCHED_DEADLINE    6

```





## 5.2.4 sched\_class连接调度器框架和调度器



## 六、CFS

```

1 struct cfs_rq {
2     struct load_weight load;
3     unsigned int nr_running;
4     unsigned int h_nr_running; /* SCHED_{NORMAL,BATCH,IDLE} */
5     unsigned int idle_h_nr_running; /* SCHED_IDLE */
6
7     u64 exec_clock;
8     u64 min_vruntime;
9
10    struct rb_root_cached tasks_timeline; /**/
11
12    /*
13     * 'curr' points to currently running entity on this cfs_rq.
14     * It is set to NULL otherwise (i.e when none are currently running).
15     */
16    struct sched_entity *curr;
17    struct sched_entity *next;
18    struct sched_entity *last;
19    struct sched_entity *skip;
20
21
22
23 #ifdef CONFIG_SMP
24    /*

```

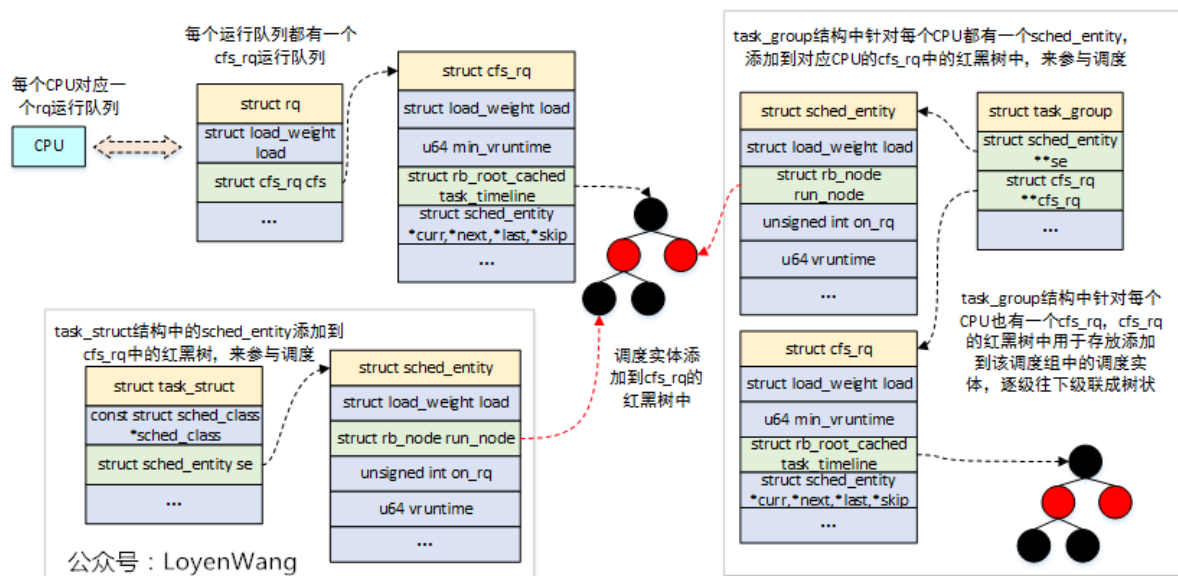
```

25     * CFS load tracking
26     */
27     struct sched_avg    avg;
28
29     struct {
30         raw_spinlock_t    lock ____cacheline_aligned;
31         int                nr;
32         unsigned long      load_avg;
33         unsigned long      util_avg;
34         unsigned long      runnable_avg;
35     } removed;
36 #endif /* CONFIG_SMP */
37
38 }

```

进程调度过程分为两部分，一是对进程信息进行修改，主要是修改和调度相关的信息，比如进程的运行时间，睡眠时间，进程的状态，cpu的负荷等等，二是进程的切换。和进程调度相关的所有函数中，只有schedule函数是用来进行进程切换的，其他函数都是用来修改进程的调度信息。

rq、cfs\_rq、task\_group、task\_struct之间的关系：




## 6.1 进程优先级

进程的优先级和调度关系密切，计算进程的虚拟运行时间要用到优先级，优先级决定进程权重，权重决定进程虚拟时间的增加速度，最终决定进程可运行时间的长短。权重越大的进程可以执行的时间越长。

### 6.1.1 优先级

在用户空间可以通过nice命令设置进程的静态优先级，这在内部会调用nice系统调用。进程的nice值在-20和+19之间（包含）。值越低，表明优先级越高。



内核使用一个简单些的数值范围，从0到139（包含），用来表示内部优先级。同样是值越低，优先级越高。从0到99的范围专供实时进程使用。nice值[-20, +19]映射到范围100到139。**实时进程的优先级总是比普通进程更高。**

```

1  #define MAX_NICE    19
2  #define MIN_NICE    -20
3  #define NICE_WIDTH  (MAX_NICE - MIN_NICE + 1)
4
5  #define MAX_USER_RT_PRIO    100
6  #define MAX_RT_PRIO        MAX_USER_RT_PRIO
7
8  #define MAX_PRIO            (MAX_RT_PRIO + NICE_WIDTH)
9  #define DEFAULT_PRIO        (MAX_RT_PRIO + NICE_WIDTH / 2)

```

```

1  static int effective_prio(struct task_struct *p)
2  {
3      p->normal_prio = normal_prio(p);
4      /*如果是实时进程或已经提高到实时优先级，则保持优先级不变。否则，返回普通优先级：*/
5      if (!rt_prio(p->prio))
6          return p->normal_prio;
7      return p->prio;
8  }

```

该函数用于设置进程的优先级，该函数设计的有一定技巧性，函数的返回值是用来设置进程的活动的优先级，但是在函数体中也把进程的普通优先级设置了。

假定我们在处理普通进程，不涉及实时调度。在这种情况下，normal\_prio只是返回静态优先级。结果很简单：所有3个优先级都是同一个值，即静态优先级！

```

1  static inline int normal_prio(struct task_struct *p) /* 获取普通优先级 */
2  {
3      int prio;
4
5      if (task_has_dl_policy(p)) /* 判断当前进程是否空闲进程，是则设置进程的普通优先级-1*/
6          prio = MAX_DL_PRIO-1;
7      else if (task_has_rt_policy(p)) /* 判断是否是实时进程，是则设置实时进程普通优先级0-99（越小优先级越高）*/
8          prio = MAX_RT_PRIO-1 - p->rt_priority;
9      else
10         prio = __normal_prio(p); /* 普通进程的普通优先级等于其静态优先级 */
11     return prio;
12 }

```

其中，第8行，看到这块减去了`p->rt_priority`，比较奇怪，这是因为实时进程描述符的`rt_priority`成员中事先存放了它自己的优先级（数字也是0-99，但在这里数字越大，优先级越高），因此往`p->prio`中倒换的时候，需要处理一下，`MAX_RT_PRIO`值为100，因此`MAX_RT_PRIO-1- (0, 99)`就倒换成了（99，0），这仅仅是个小技巧。

### 6.1.2 权重

进程的重要性不仅是由优先级指定的，而且还需要考虑保存在`task_struct->se.load`的负荷权重。

`include/linux/sched.h` 中定义了权重的结构体：

```
1 struct load_weight {
2     unsigned long    weight;
3     u32              inv_weight;
4 };
```

```
1 const int sched_prio_to_weight[40] = {
2     /* -20 */      88761,      71755,      56483,      46273,      36291,
3     /* -15 */      29154,      23254,      18705,      14949,      11916,
4     /* -10 */      9548,       7620,       6100,       4904,       3906,
5     /* -5  */      3121,       2501,       1991,       1586,       1277,
6     /*  0  */      1024,       820,        655,        526,        423,
7     /*  5  */       335,       272,        215,        172,        137,
8     /* 10  */       110,       87,         70,         56,         45,
9     /* 15  */       36,       29,         23,         18,         15,
10 };
```

进程每降低一个nice值，则多获得10%的CPU时间，每升高一个nice值，则放弃10%的CPU时间。对内核使用的范围[0, 39]中的每个nice级别，该数组中都有一个对应项。各数组之间的乘数因子是1.25。

### 设置权重

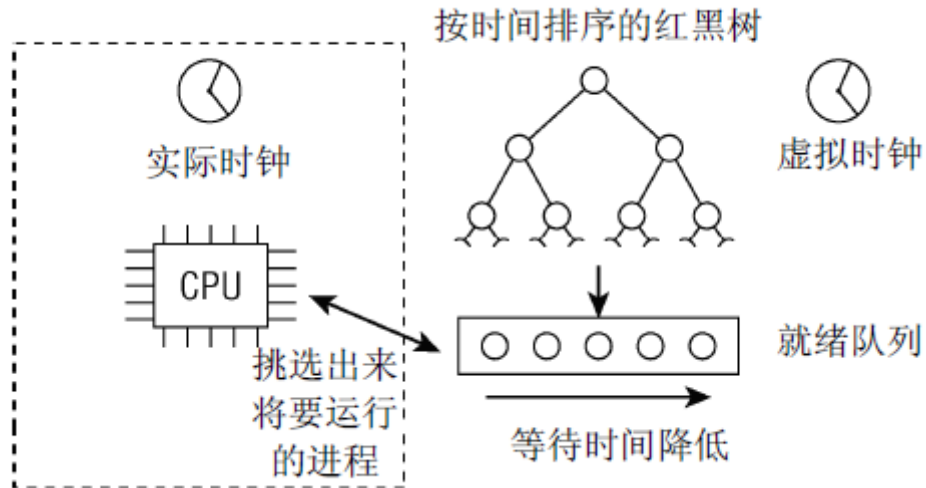
```
1 static void set_load_weight(struct task_struct *p, bool update_load)
2 {
3     int prio = p->static_prio - MAX_RT_PRIO;
4     struct load_weight *load = &p->se.load; /* 权重保存在task_struct的
5 se.load中 */
6     /*SCHED_IDLE进程得到的权重最小：*/
7     if (task_has_idle_policy(p)) {
8         load->weight = scale_load(WEIGHT_IDLEPRIO);
9         load->inv_weight = WMULT_IDLEPRIO;
10        return;
11    }
12
13    /*
14     * SCHED_OTHER tasks have to update their load when changing their
15     * weight
16     */
17    if (update_load && p->sched_class == &fair_sched_class) {
18        reweight_task(p, prio);
19    } else {
20        load->weight = scale_load(sched_prio_to_weight[prio]);
```

```

21     load->inv_weight = sched_prio_to_wmult[prio];
22 }
23 }
24

```

## 虚拟时间



所有的可运行进程都按时间在一个红黑树中排序，所谓时间即其等待时间。等待CPU时间最长的进程是最左侧的项，调度器下一次会考虑该进程。等待时间稍短的进程在该树上从左至右排序。

完全公平调度算法依赖于虚拟时钟，用以度量等待进程在完全公平系统中所能得到的CPU时间。

所有与虚拟时钟有关的计算都在update\_curr中执行，该函数在系统中各个不同地方调用，包括周期性调度器之内

```

1  /*
2   * Update the current task's runtime statistics.
3   */
4  static void update_curr(struct cfs_rq *cfs_rq)
5  {
6      struct sched_entity *curr = cfs_rq->curr;
7      u64 now = rq_clock_task(rq_of(cfs_rq)); /* 从就绪队列rq的clock_task成员
8      中获取当前时间 */
9      u64 delta_exec;
10
11      if (unlikely(!curr))
12          return;
13
14      delta_exec = now - curr->exec_start; /*当前时间减去进程上次时钟中断tick中
15      开始时间得到进程运行的时间间隔*/
16      if (unlikely((s64)delta_exec <= 0))
17          return;
18
19      curr->exec_start = now; /* 当前时间赋值给进程新的开始时间 */
20
21      schedstat_set(curr->statistics.exec_max,
22                    max(delta_exec, curr->statistics.exec_max));
23

```

```

22  /*将进程运行的时间间隔delta_exec累加到调度实体的sum_exec_runtime成员中，该成员代
    表进程到目前为止运行了多长时间*/
23  curr->sum_exec_runtime += delta_exec;
24  schedstat_add(cfs_rq->exec_clock, delta_exec); /*将进程运行的时间间隔
    delta_exec也累加到公平调度就绪队列cfs_rq的exec_clock成员中*/
25
26  /*calc_delta_fair函数很关键，它将进程执行的真实运行时间转换成虚拟运行时间，然后累加
    到调度实体的vruntime域中*/
27  curr->vruntime += calc_delta_fair(delta_exec, curr);
28  update_min_vruntime(cfs_rq); /*更新cfs_rq队列中的最小虚拟运行时间
    min_vruntime，该时间是就绪队列中所有进程包括当前进程的已运行的最小虚拟时间，只能单调递增
    */
29
30  if (entity_is_task(curr)) {
31      struct task_struct *curtask = task_of(curr);
32
33      trace_sched_stat_runtime(curtask, delta_exec, curr->vruntime);
34      cgroup_account_cputime(curtask, delta_exec);
35      account_group_exec_runtime(curtask, delta_exec);
36  }
37
38  account_cfs_rq_runtime(cfs_rq, delta_exec);
39  }
40

```

每个cfs\_rq队列均有一个min\_vruntime成员，装的是就绪队列中所有进程包括当前进程已运行的虚拟时间中最小的那个时间。update\_min\_vruntime用于更新该时间。

**队列中的min\_vruntime成员非常重要**，用于在睡眠进程被唤醒后以及新进程被创建好时，进行虚拟时间补偿或者惩罚

```

1  static void update_min_vruntime(struct cfs_rq *cfs_rq)
2  {
3      struct sched_entity *curr = cfs_rq->curr;
4      struct rb_node *leftmost = rb_first_cached(&cfs_rq->tasks_timeline);
5
6      u64 vruntime = cfs_rq->min_vruntime;
7
8      if (curr) {
9          if (curr->on_rq)
10             vruntime = curr->vruntime;
11         else
12             curr = NULL;
13     }
14
15     if (leftmost) { /* non-empty tree */ /*就绪队列中有下一个要被调度的进程，则进入下一个调度实体*/
16         struct sched_entity *se;
17         se = rb_entry(leftmost, struct sched_entity, run_node);
18
19         /*从当前进程和下一个被调度进程中，选择最小的已运行虚拟时间，保存到vruntime中*/
20         if (!curr)
21             vruntime = se->vruntime;
22         else
23             vruntime = min_vruntime(vruntime, se->vruntime);

```

```

24     }
25
26     /*从当前队列的min_vruntime域和vruntime变量中，选最大的保存到队列的min_vruntime域
    中，完成更新*/
27     /* ensure we never gain time by being placed backwards. */
28     cfs_rq->min_vruntime = max_vruntime(cfs_rq->min_vruntime, vruntime);
29 #ifndef CONFIG_64BIT
30     smp_wmb();
31     cfs_rq->min_vruntime_copy = cfs_rq->min_vruntime;
32 #endif
33 }

```

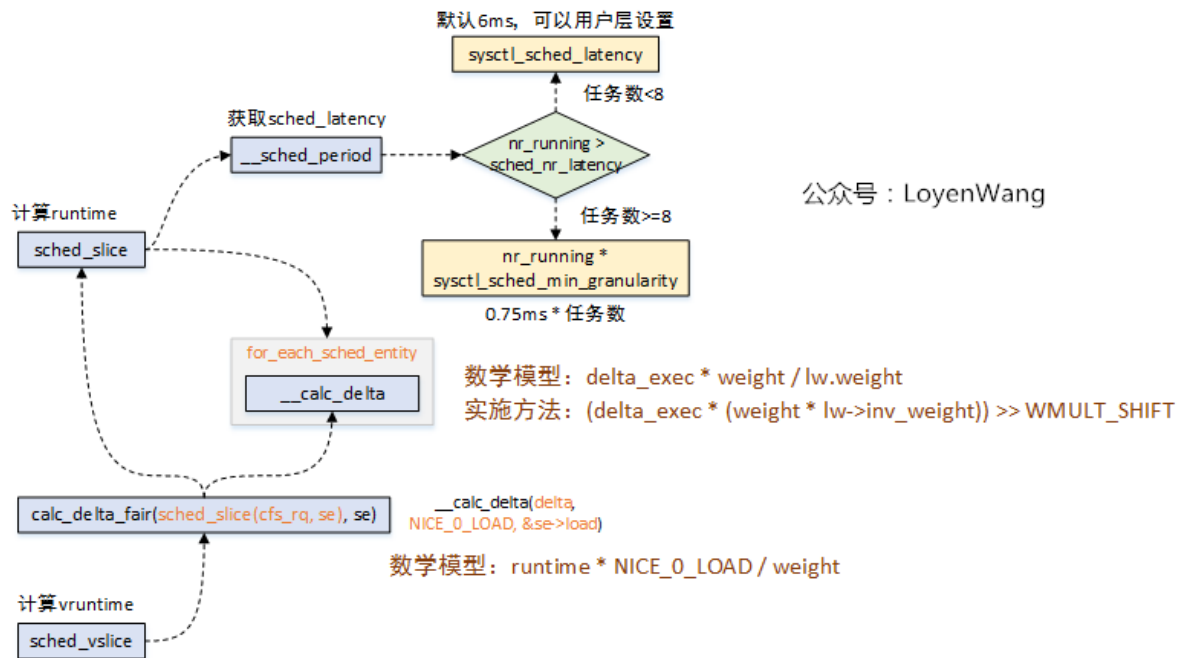
```

1  static inline u64 calc_delta_fair(u64 delta, struct sched_entity *se)
2  {
3      if (unlikely(se->load.weight != NICE_0_LOAD))
4          delta = __calc_delta(delta, NICE_0_LOAD, &se->load);
5
6      return delta;
7  }
8
9  static u64 __calc_delta(u64 delta_exec, unsigned long weight, struct
    load_weight *lw)
10 {
11     u64 fact = scale_load_down(weight);
12     int shift = WMULT_SHIFT;
13
14     __update_inv_weight(lw);
15
16     if (unlikely(fact >> 32)) {
17         while (fact >> 32) {
18             fact >>= 1;
19             shift--;
20         }
21     }
22
23     fact = mul_u32_u32(fact, lw->inv_weight);
24
25     while (fact >> 32) {
26         fact >>= 1;
27         shift--;
28     }
29
30     return mul_u64_u32_shr(delta_exec, fact, shift);
31 }

```

## 虚拟时间到底怎么回事？

sched\_vslice 函数计算虚拟时间



## 负载

内核中计算CPU负载的方法是PELT(Per-Entity Load Tracing), 不仅考虑进程权重, 而且跟踪每个调度实体的负载情况。

sched\_entity结构中有一个struct sched\_avg用于描述进程的负载

```

1 struct sched_avg {
2     u64         last_update_time;
3     u64         load_sum;
4     u64         runnable_sum;
5     u32         util_sum;
6     u32         period_contrib;
7     unsigned long load_avg;
8     unsigned long runnable_avg;
9     unsigned long util_avg;
10    struct util_est util_est;
11 } ____cacheline_aligned;

```

### 6.1.3 选择下一个进程

```

1 /*
2  * Pick up the highest-prio task:
3  */
4 static inline struct task_struct *
5 pick_next_task(struct rq *rq, struct task_struct *prev, struct rq_flags
6 *rf)
7 {
8     const struct sched_class *class;

```



```

8     struct task_struct *p;
9
10    if (likely(prev->sched_class == &fair_sched_class &&
11            rq->nr_running == rq->cfs.h_nr_running)) {
12
13        p = pick_next_task_fair(rq, prev, rf);
14        if (unlikely(p == RETRY_TASK))
15            goto restart;
16
17        /* Assumes fair_sched_class->next == idle_sched_class */
18        if (!p) {
19            put_prev_task(rq, prev);
20            p = pick_next_task_idle(rq);
21        }
22
23        return p;
24    }
25
26restart:
27    put_prev_task_balance(rq, prev, rf);
28
29    for_each_class(class) {
30        p = class->pick_next_task(rq);
31        if (p)
32            return p;
33    }
34
35    /* The idle class should always have a runnable task: */
36    BUG();
37 }
38
39
40
41
42
43 /*主要是pick_next_task_fair函数*/
44 struct task_struct *
45 pick_next_task_fair(struct rq *rq, struct task_struct *prev, struct
46 rq_flags *rf)
47 {
48     struct cfs_rq *cfs_rq = &rq->cfs;
49     struct sched_entity *se;
50     struct task_struct *p;
51     int new_tasks;
52
53again:
54     if (!sched_fair_runnable(rq))
55         goto idle;
56
57#ifdef CONFIG_FAIR_GROUP_SCHED
58     if (!prev || prev->sched_class != &fair_sched_class)
59         goto simple;
60
61     do {
62         struct sched_entity *curr = cfs_rq->curr;
63
64         if (curr) {
65             if (curr->on_rq)

```

```

65         update_curr(cfs_rq);
66     else
67         curr = NULL;
68
69     if (unlikely(check_cfs_rq_runtime(cfs_rq))) {
70         cfs_rq = &rq->cfs;
71
72         if (!cfs_rq->nr_running)
73             goto idle;
74
75         goto simple;
76     }
77 }
78
79 se = pick_next_entity(cfs_rq, curr);
80 cfs_rq = group_cfs_rq(se);
81 } while (cfs_rq); /*对所有的调度组进行遍历，从中选择下一个可调度的进程，而不只局
    限在当前队列的当前组*/
82
83 p = task_of(se);
84
85 if (prev != p) {
86     struct sched_entity *pse = &prev->se;
87
88     while (!(cfs_rq = is_same_group(se, pse))) {
89         int se_depth = se->depth;
90         int pse_depth = pse->depth;
91
92         if (se_depth <= pse_depth) {
93             put_prev_entity(cfs_rq_of(pse), pse);
94             pse = parent_entity(pse);
95         }
96         if (se_depth >= pse_depth) {
97             set_next_entity(cfs_rq_of(se), se);
98             se = parent_entity(se);
99         }
100     }
101
102     put_prev_entity(cfs_rq, pse);
103     set_next_entity(cfs_rq, se);
104 }
105
106 goto done;
107 simple:
108 #endif
109 if (prev)
110     put_prev_task(rq, prev);
111
112 do {
113     se = pick_next_entity(cfs_rq, NULL);
114     set_next_entity(cfs_rq, se);
115     cfs_rq = group_cfs_rq(se);
116 } while (cfs_rq);
117
118 p = task_of(se);
119
120 done: __maybe_unused;
121 #ifdef CONFIG_SMP

```

```

122     /*
123      * Move the next running task to the front of
124      * the list, so our cfs_tasks list becomes MRU
125      * one.
126      */
127     list_move(&p->se.group_node, &rq->cfs_tasks);
128 #endif
129
130     if (hrtick_enabled(rq))
131         hrtick_start_fair(rq, p);
132
133     update_misfit_status(p, rq);
134
135     return p;
136
137 idle:
138     if (!rf)
139         return NULL;
140
141     new_tasks = newidle_balance(rq, rf);
142
143     /*
144      * Because newidle_balance() releases (and re-acquires) rq->lock, it is
145      * possible for any higher priority task to appear. In that case we
146      * must re-start the pick_next_entity() loop.
147      */
148     if (new_tasks < 0)
149         return RETRY_TASK;
150
151     if (new_tasks > 0)
152         goto again;
153
154     /*
155      * rq is about to be idle, check if we need to update the
156      * lost_idle_time of clock_pelt
157      */
158     update_idle_rq_clock_pelt(rq);
159
160     return NULL;
161 }

```

## 6.1.4 就绪队列的入队和出队

### enqueue\_task\_fair()函数

CFS的enqueue\_task钩子函数是 `enqueue_task_fair()` 函数：

```

1  /*
2   * The enqueue_task method is called before nr_running is
3   * increased. Here we update the fair scheduling stats and
4   * then put the task into the rbtree:
5   */
6  /**nr_running是cfs_rq结构体中的成员，计数所有就绪的进程数包括cfs_rq中以及正在运行的进
   程**/
7  static void
8  enqueue_task_fair(struct rq *rq, struct task_struct *p, int flags)

```

```

9 {
10     struct cfs_rq *cfs_rq;
11     struct sched_entity *se = &p->se;          /*获取进程p的调度实体*/
12     int idle_h_nr_running = task_has_idle_policy(p);    /*判断进程 p->policy
== SCHED_IDLE, 有何作用? /
13
14     /*
15      * The code below (indirectly) updates schedutil which looks at
16      * the cfs_rq utilization to select a frequency.
17      * Let's add the task's estimated utilization to the cfs_rq's
18      * estimated utilization, before we update schedutil.
19      */
20     util_est_enqueue(&rq->cfs, p); /*cfs的负载估算, 占用cpu的算力负载估算*/
21
22     /*
23      * If in_iowait is set, the code below may not trigger any cpufreq
24      * utilization updates, so do it here explicitly with the IOWAIT flag
25      * passed.
26      */
27     if (p->in_iowait)
28         cpufreq_update_util(rq, SCHED_CPUFREQ_IOWAIT);
29
30     /*这段代码在做什么? */
31     for_each_sched_entity(se) {
32         if (se->on_rq) /*判断进程是否已经在队列里, on_rq为1则不需要再加入队列了, 已
经存在队列里*/
33             break;
34         cfs_rq = cfs_rq_of(se);
35         enqueue_entity(cfs_rq, se, flags); /*将调度实体加入队列*/
36
37         cfs_rq->h_nr_running++; /*计数增加*/
38         cfs_rq->idle_h_nr_running += idle_h_nr_running;
39
40         /* end evaluation on encountering a throttled cfs_rq */
41         if (cfs_rq_throttled(cfs_rq))
42             goto enqueue_throttle;
43
44         flags = ENQUEUE_WAKEUP;
45     }
46
47     /*为什么需要两次循环, 操作不同在什么地方呢? */
48     for_each_sched_entity(se) {
49         cfs_rq = cfs_rq_of(se);
50
51         update_load_avg(cfs_rq, se, UPDATE_TG);
52         se_update_runnable(se);
53         update_cfs_group(se);
54
55         cfs_rq->h_nr_running++;
56         cfs_rq->idle_h_nr_running += idle_h_nr_running;
57
58         /* end evaluation on encountering a throttled cfs_rq */
59         if (cfs_rq_throttled(cfs_rq))
60             goto enqueue_throttle;
61
62         /*
63          * One parent has been throttled and cfs_rq removed from the
64          * list. Add it back to not break the leaf list.

```

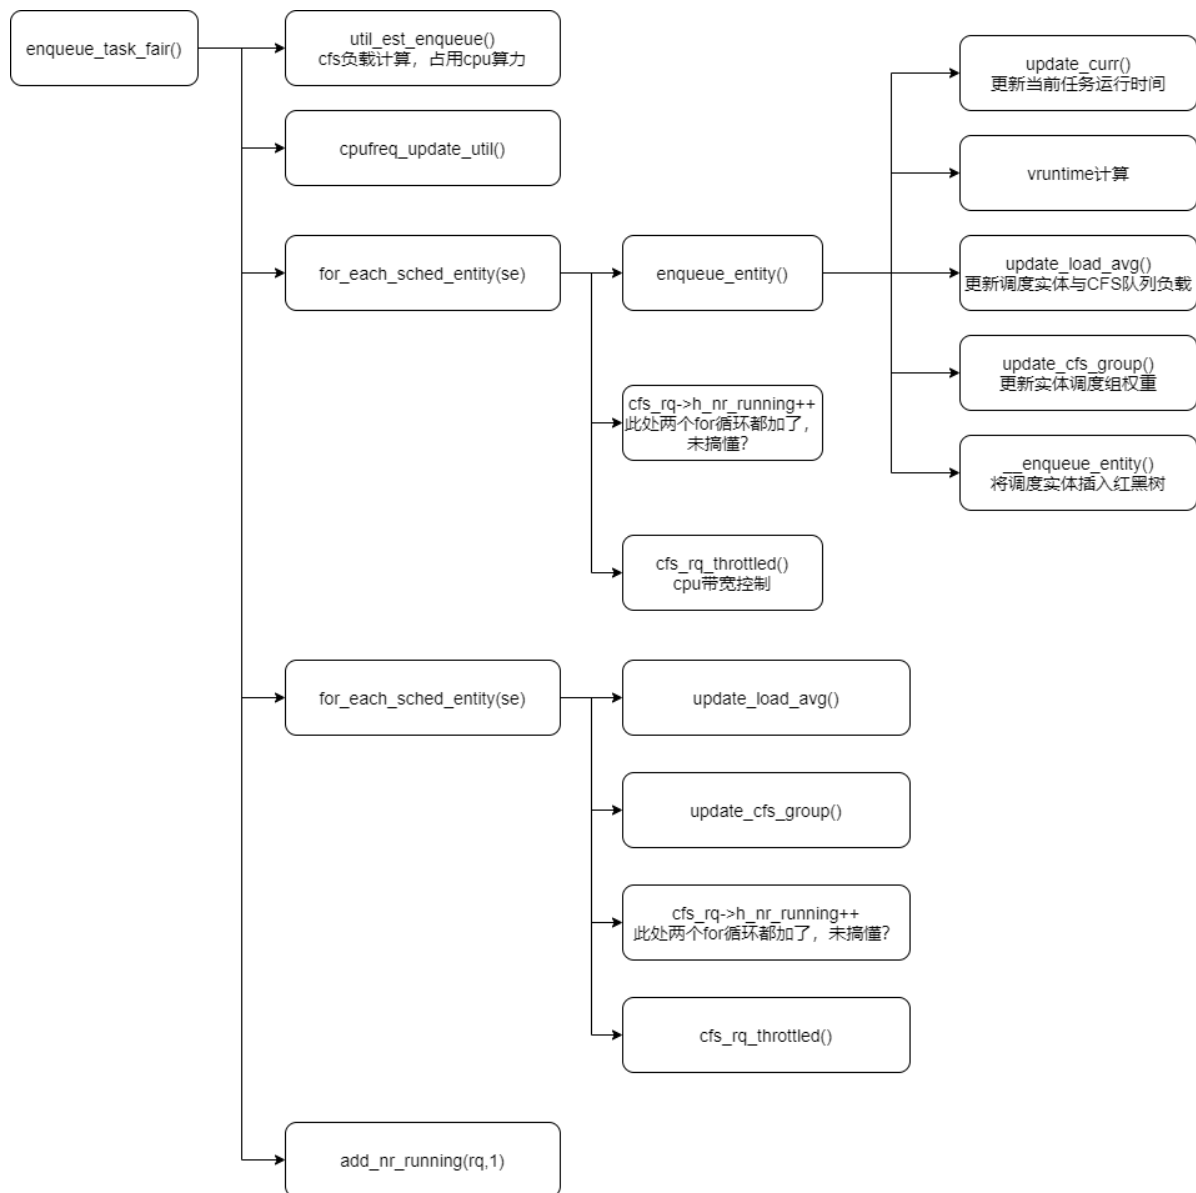
```

65         */
66         if (throttled_hierarchy(cfs_rq))
67             list_add_leaf_cfs_rq(cfs_rq);
68     }
69
70     /* At this point se is NULL and we are at root level*/
71     add_nr_running(rq, 1);
72
73     /*
74      * Since new tasks are assigned an initial util_avg equal to
75      * half of the spare capacity of their CPU, tiny tasks have the
76      * ability to cross the overutilized threshold, which will
77      * result in the load balancer ruining all the task placement
78      * done by EAS. As a way to mitigate that effect, do not account
79      * for the first enqueue operation of new tasks during the
80      * overutilized flag detection.
81      *
82      * A better way of solving this problem would be to wait for
83      * the PELT signals of tasks to converge before taking them
84      * into account, but that is not straightforward to implement,
85      * and the following generally works well enough in practice.
86      */
87     if (flags & ENQUEUE_WAKEUP)
88         update_overutilized_status(rq);
89
90 enqueue_throttle:
91     if (cfs_bandwidth_used()) {
92         /*
93          * When bandwidth control is enabled; the cfs_rq_throttled()
94          * breaks in the above iteration can result in incomplete
95          * leaf list maintenance, resulting in triggering the assertion
96          * below.
97          */
98         for_each_sched_entity(se) {
99             cfs_rq = cfs_rq_of(se);
100
101             if (list_add_leaf_cfs_rq(cfs_rq))
102                 break;
103         }
104     }
105
106     assert_list_leaf_cfs_rq(rq);
107
108     hrtick_update(rq);
109 }

```

enqueue\_task\_fair主要职责：

- 1) 更新运行时的数据，比如负载、权重、组调度的占比等等；
- 2) 将sched\_entity插入红黑树；



将调度实体入队红黑树。

```

1  /*
2   * Enqueue an entity into the rb-tree:
3   */
4  static void __enqueue_entity(struct cfs_rq *cfs_rq, struct sched_entity *se)
5  {
6      struct rb_node **link = &cfs_rq->tasks_timeline.rb_root.rb_node; /*获取
7      就绪队列中红黑树的根节点*/
8      struct rb_node *parent = NULL; /* 用于指向树根*/
9      struct sched_entity *entry;
10     bool leftmost = true;
11
12     /*
13      * Find the right place in the rbtree:
14      */
15     while (*link) {
16         parent = *link;
17         entry = rb_entry(parent, struct sched_entity, run_node); /*获得树根节
18         点的调度实体*/
19     }
20     rb_insert_node(link, entry);
21 }

```

```

17
18     /*比较要入队的实体中的已运行虚拟时间和树根实体中的该信息，如果前者小的话，就要插入
    到树的左子树上（link指向树根的左孩子，再次进入循环，类似于递归），否则就要插入到树的右子树
    上（同上）。这块就将进程的调度策略展现的淋漓尽致：根据进程已运行的虚拟时间来决定进程的调度，
    红黑树的左子树比右子树要先被调度，已运行的虚拟时间越小的进程越在树的左侧*/
19     if (entity_before(se, entry)) {
20         link = &parent->rb_left;
21     } else {
22         link = &parent->rb_right;
23         leftmost = false;
24     }
25 }
26
27 rb_link_node(&se->run_node, parent, link); /*红黑树重新着色*/
28 rb_insert_color_cached(&se->run_node,
29                        &cfs_rq->tasks_timeline, leftmost);
30 }

```

### dequeue\_task\_fair()函数

```

1  /*
2   * The dequeue_task method is called before nr_running is
3   * decreased. We remove the task from the rbtree and
4   * update the fair scheduling stats:
5   */
6  static void dequeue_task_fair(struct rq *rq, struct task_struct *p, int
    flags)
7  {
8      struct cfs_rq *cfs_rq;
9      struct sched_entity *se = &p->se;
10     int task_sleep = flags & DEQUEUE_SLEEP;
11     int idle_h_nr_running = task_has_idle_policy(p);
12     bool was_sched_idle = sched_idle_rq(rq);
13
14     for_each_sched_entity(se) {
15         cfs_rq = cfs_rq_of(se);
16         dequeue_entity(cfs_rq, se, flags);
17
18         cfs_rq->h_nr_running--;
19         cfs_rq->idle_h_nr_running -= idle_h_nr_running;
20
21         /* end evaluation on encountering a throttled cfs_rq */
22         if (cfs_rq_throttled(cfs_rq))
23             goto dequeue_throttle;
24
25         /* Don't dequeue parent if it has other entities besides us */
26         if (cfs_rq->load.weight) {
27             /* Avoid re-evaluating load for this entity: */
28             se = parent_entity(se);
29             /*
30              * Bias pick_next to pick a task from this cfs_rq, as
31              * p is sleeping when it is within its sched_slice.
32              */
33             if (task_sleep && se && !throttled_hierarchy(cfs_rq))
34                 set_next_buddy(se);

```

```

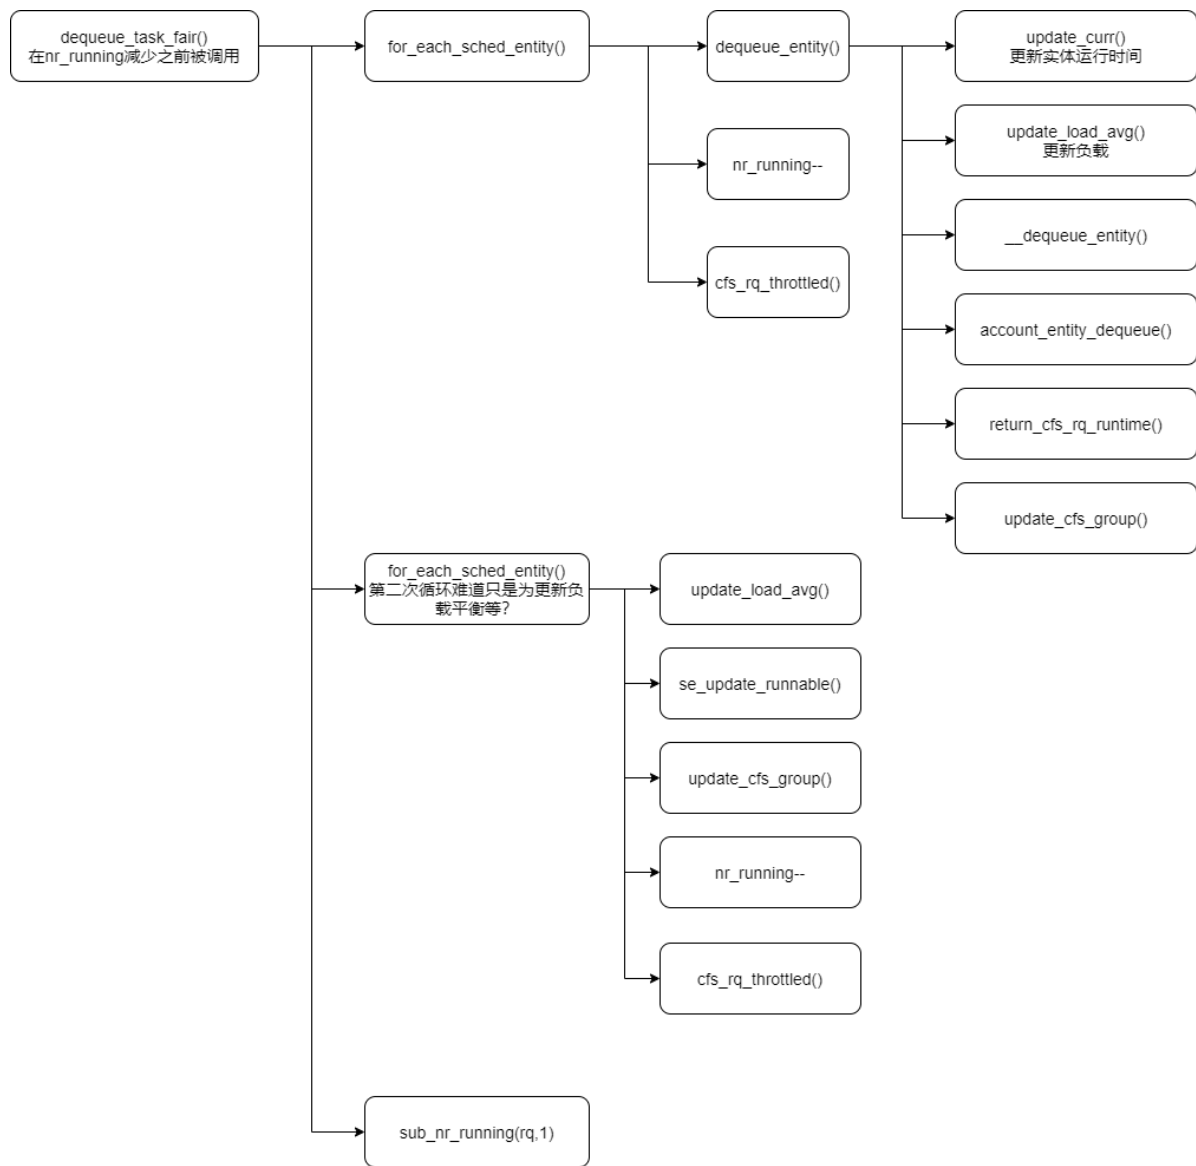
35         break;
36     }
37     flags |= DEQUEUE_SLEEP;
38 }
39
40 for_each_sched_entity(se) {
41     cfs_rq = cfs_rq_of(se);
42
43     update_load_avg(cfs_rq, se, UPDATE_TG);
44     se_update_runnable(se);
45     update_cfs_group(se);
46
47     cfs_rq->h_nr_running--;
48     cfs_rq->idle_h_nr_running -= idle_h_nr_running;
49
50     /* end evaluation on encountering a throttled cfs_rq */
51     if (cfs_rq_throttled(cfs_rq))
52         goto dequeue_throttle;
53
54 }
55
56 /* At this point se is NULL and we are at root level*/
57 sub_nr_running(rq, 1);
58
59 /* balance early to pull high priority tasks */
60 if (unlikely(!was_sched_idle && sched_idle_rq(rq)))
61     rq->next_balance = jiffies;
62
63 dequeue_throttle:
64     util_est_dequeue(&rq->cfs, p, task_sleep);
65     hrtick_update(rq);
66 }

```

dequeue\_task\_fair的主要工作内容和enqueue其实类型：

- 1) 更新运行时间、负载等；
- 2) 将实体移出红黑树队列；





## 出队列

```

1 static void __dequeue_entity(struct cfs_rq *cfs_rq, struct sched_entity *se)
2 {
3     rb_erase_cached(&se->run_node, &cfs_rq->tasks_timeline);
4 }
5
6 static inline void rb_erase_cached(struct rb_node *node,
7     struct rb_root_cached *root)
8 {
9     if (root->rb_leftmost == node) /*判断要出队的实体是不是红黑树最左侧的孩子
10         (rb_leftmost所指向的)*/
11         root->rb_leftmost = rb_next(node); /*是最左子树的话需要找出下一个*/
12     rb_erase(node, &root->rb_root);
13 }

```

### 6.1.5睡眠进程被唤醒后抢占当前进程

当某个资源空出来后，等待该资源的进程就会被唤醒，唤醒后也许就要抢占当前进程。

该函数会唤醒睡眠中的指定p的进程。

```
1 static int
2 try_to_wake_up(struct task_struct *p, unsigned int state, int wake_flags)
3 {
4     [...]
5 }
```

唤醒一个刚被创建的进程

```
1 /*
2  * wake_up_new_task - wake up a newly created task for the first time.
3  *
4  * This function will do some initial scheduler statistics housekeeping
5  * that must be done for every newly created context, then puts the task
6  * on the runqueue and wakes it.
7  */
8 void wake_up_new_task(struct task_struct *p)
9 {
10     [...]
11 }
```

检查唤醒进程是否能抢占当前进程.

```
1 /*
2  * Preempt the current task with a newly woken task if needed:
3  */
4 static void check_preempt_wakeup(struct rq *rq, struct task_struct *p, int
wake_flags)
5 {
6     [...]
7 }
```

### 6.1.6 fork的处理

该函数在do\_fork-->copy\_process函数中调用，用来设置新创建进程的虚拟时间信息。

```
1 /*
2  * called on fork with the child task as argument from the parent's context
3  * - child not yet on the tasklist
4  * - preemption disabled
5  */
6 static void task_fork_fair(struct task_struct *p)
```

```

7  {
8      struct cfs_rq *cfs_rq;
9      struct sched_entity *se = &p->se, *curr;
10     struct rq *rq = this_rq();
11     struct rq_flags rf;
12
13     rq_lock(rq, &rf);
14     update_rq_clock(rq);
15
16     cfs_rq = task_cfs_rq(current);
17     curr = cfs_rq->curr;
18     if (curr) {
19         update_curr(cfs_rq);
20         se->vruntime = curr->vruntime;      /*当前进程（父进程）的虚拟运行时间拷贝
给新进程（子进程）*/
21     }
22     place_entity(cfs_rq, se, 1);            /*完成新进程的“时间片”计算以及虚拟时间
惩罚，之后将新进程加入红黑树中*/
23
24     /*如果设置了子进程先于父进程运行的标志并且当前进程不为空且当前进程已运行的虚拟时间比新
进程小，则执行if体*/
25     if (sysctl_sched_child_runs_first && curr && entity_before(curr, se)) {
26         /*交换当前进程和新进程的虚拟时间（新进程的虚拟时间变小，就排在了红黑树的左侧，当前
进程之前，下次就能被调度）*/
27         swap(curr->vruntime, se->vruntime);
28         resched_curr(rq); /*设置重新调度标志*/
29     }
30
31     se->vruntime -= cfs_rq->min_vruntime; /*给新进程的虚拟运行时间减去队列的最小虚
拟时间来做一点补偿（因为在上边的place_entity函数中给新进程的虚拟时间加了一次
min_vruntime，所以在这里要减去）*/
32     rq_unlock(rq, &rf);
33 }

```

看下place\_entity函数，该函数完成新进程的“时间片”计算和虚拟时间惩罚，并且将新进程加入就绪队列。

```

1  static void
2  place_entity(struct cfs_rq *cfs_rq, struct sched_entity *se, int initial)
3  {
4      u64 vruntime = cfs_rq->min_vruntime;
5
6      /*如果initial标志为1的话（说明当前计算的是新进程的时间），将计算出的新进程的虚拟时间
片累加到vruntime中，累加到原因是调度系统要保证先把就绪队列中的所有的进程执行一遍之后才能执
行新进程*/
7      if (initial && sched_feat(START_DEBIT))
8          vruntime += sched_vslice(cfs_rq, se);
9
10     /* sleeps up to a single latency don't count. */
11     if (!initial) {                /*如果当前计算的不是新进程（睡眠的进程），把一个延迟
周期的长度sysctl_sched_latency（6ms）赋给thresh*/
12         unsigned long thresh = sysctl_sched_latency;
13
14         /*
15          * Halve their sleep time's effect, to allow

```

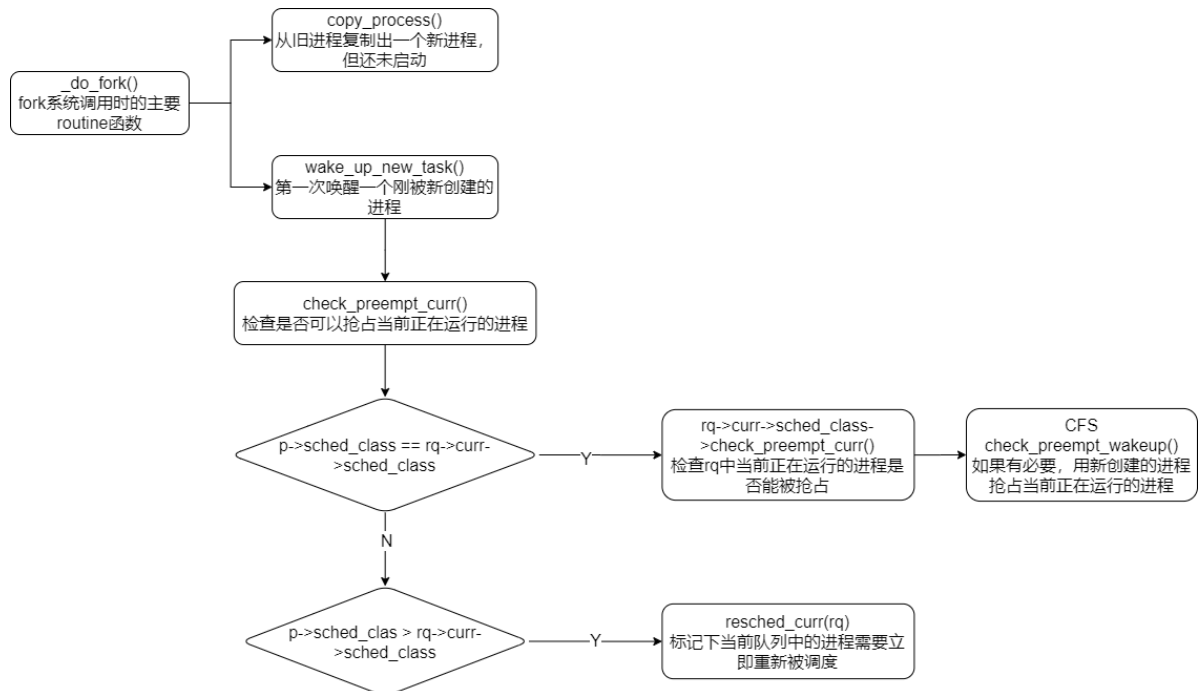
```

16      * for a gentler effect of sleepers:
17      */
18      if (sched_feat(GENTLE_FAIR_SLEEPERS))
19          thresh >= 1;          /*thresh减半*/
20
21      vruntime -= thresh; /*睡眠进程的虚拟运行时间减去减半后的thresh，因为睡眠进程
    好长时间未运行，因此要进行虚拟时间补偿，把它已运行的虚拟时间减小一点，使得它能多运行一会*/
22  }
23
24      /* ensure we never gain time by being placed backwards. */
25      se->vruntime = max_vruntime(se->vruntime, vruntime);          /*将设置
    好的虚拟时 间保存到进程调度实体的vruntime域*/
26  }

```

为什么要对新进程进行虚拟时间惩罚，其实原因只有一个，就是调度系统要保证将就绪队列中现有的进程执行一遍之后再执行新进程，那么就使新进程的  $vruntime = cfs\_rq \rightarrow min\_vruntime + \text{新进程的虚拟时间片}$ ，才能使得新进程插入到红黑树的右边，最后参与调度，不然无法保证所有进程在新进程之前执行。

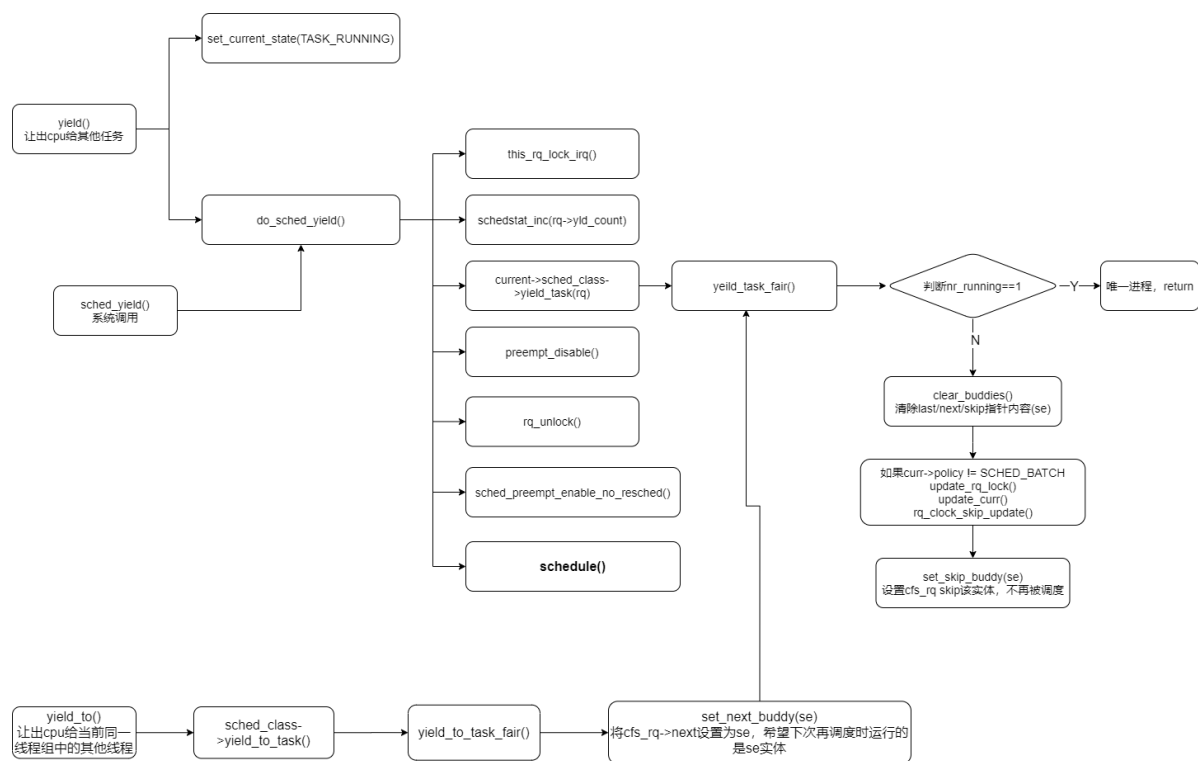
### check\_preempt\_curr()函数



### 6.1.7 让出cpu

两个yield\_task相关函数:

- `yield()`让出cpu给其他线程
- `yield_to()`让出cpu给同一线程去的其他线程（是否是同一队列？）



查看MuQSS调度器是否启用。

全称: The Multiple Queue Skiplist Scheduler

## 7.1 skiplist跳表

跳跃表使用概率均衡技术而不是使用强制性均衡，因此，对于插入和删除结点比传统上的平衡树算法更为简洁高效。

## 7.2 MuQSS简介

BFS虽然简单，但是两个问题却非常明显：

## 2. 多CPU操作全局链表的锁问题。

我们看看BFS的算法简单到何种程度：

- task插入：直接将task插入链表末尾。
- task选择：冒泡选择Virtual Deadline最小的task。【在遍历过程中会有trick，发现当前jiffies大于task的VD，就退出，这像极了Linux内核的timer处理】

最终，Con Kolivas认为：

1. 在task数量并不太大的情况下， $O(n)$ 算法没有任何问题。
2. 在CPU数量保持在16个以内时，争锁的开销可以忽略。

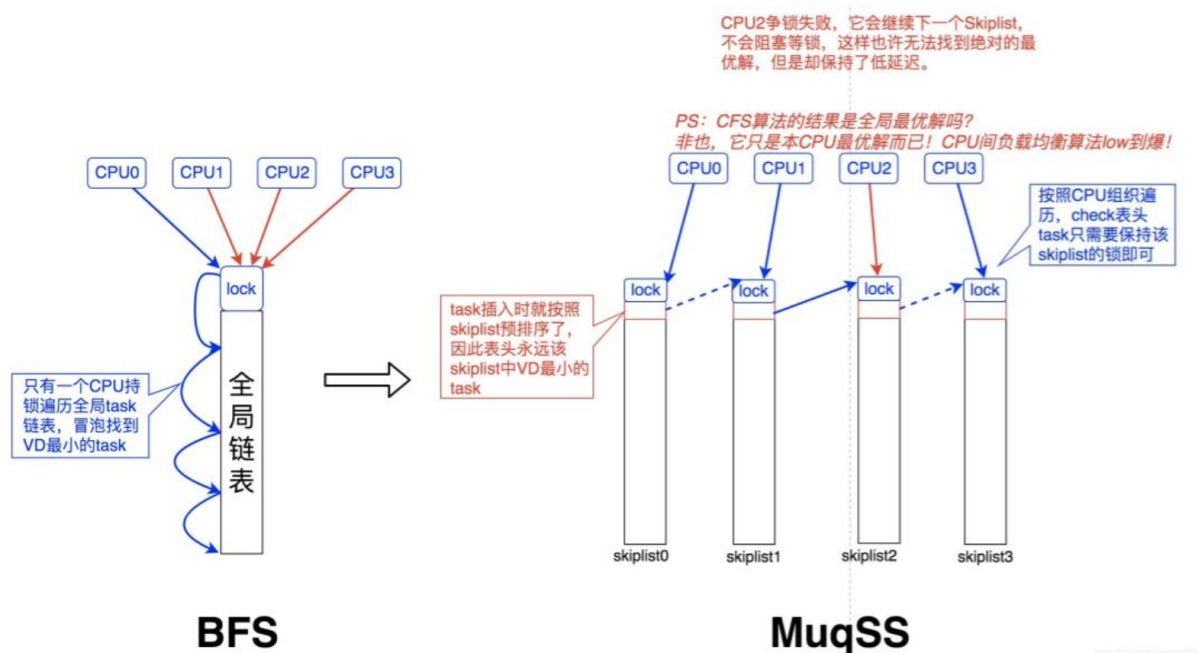
MuqSS零代价解决了BFS存在的两个问题：

1. 遍历查找的 $O(n)$ 问题。引入Skiplist数据结构替换双向链表，在 $O(\log n)$ 的插入代价下将查找的时间复杂度降为 $O(1)$ 。【关于Skiplist，可以参考我的另一篇文章：<https://blog.csdn.net/dog250/article/details/46997155>】
2. 多CPU操作全局链表的锁问题。引入每CPU链表，避免全局争锁。同时以trylock代替lock，以损失准确性为代价实现无锁操作。

Con Kolivas在 *保持简单* 这个约束下设计了MuqSS，其要点是：

- Skiplist的作用类似主线Linux内核CFS中的红黑树，但比红黑树简单得多。
- 选择task的算法遍历所有CPU的Skiplist表头，选择当前全局最优task。
- 锁粒度细化到每个CPU的Skiplist。
- 遍历过程针对每CPU锁采用trylock，失败则继续下一个CPU，实现无锁化。

时间复杂度同样都是 $O(n)$ ，但MuqSS的n指的是CPU数量而非task数量。



## 7.3 MuQSS关键因子

proc参数	默认值	含义	
iso_cpu	70	该值设置了无特权的SCHED_ISO进程可以以实施优先级运行的cpu百分占比，即在整个系统(即所有cpu)上滚动5秒的平均cpu百分比。 SCHED_ISO在linux-5.9.10中保留了，并未实现。	MuQSS独有
kexec_load_disabled	0	ROM/Flash boot loader	
rr_interval	6	MuQSS独有；该值是多少cpu调度单元可以运行的最小时间长度。增加该值可以提高计算密集型任务的吞吐量，但会增加延迟；同样，减少该值，牺牲吞吐量，降低了平均和最大延迟。 该值是ms级别，可设置范围为1-1000，一般默认值是根据调度器初始化时可用的cpu数量来决定，一般最小为6；	MuQSS独有
sched_energy_aware	1	softlockup threshold，是看门狗threshold的两倍大。如果将该值设置为0，则会关闭lockup探测。	
yield_type	1	该值决定了sched_yield函数调用时会怎么表现 0：不放弃cpu 1：只放弃cpu给更高优先级的进程 2：耗尽时间片并重新计算deadline	MuQSS独有

## 7.4 MuQSS的源码实现

### 7.4.1 skip\_list分析

MuQSS的skip\_list主要增加了两个文件 `include/linux/skip_list.h` 以及 `kernel/skip_list.c`

skip\_list.h中：

```
1  typedef u64 keyType;
2  typedef void *valueType;
3  typedef struct nodeStructure skiplist_node;
4
5  struct nodeStructure {
6      int level; /* Levels in this structure */
7      keyType key;
8      valueType value;
9      skiplist_node *next[8]; /*这里一共是8个，和后面的MaxNumberOfLevels对应么?*/
10     skiplist_node *prev[8];
11 }; //定义跳表节点
12
13 typedef struct listStructure {
14     int entries; /*记录元素个数，每次插入加1，每次删除减1*/
15     int level; /* Maximum level of the list
16                (1 more than the number of levels in the list) */
17     skiplist_node *header; /* pointer to header */
```

```
18 } skiplist;
```

skiplist\_node节点结构体中有个level，skiplist表结构体中也有个level，这两个level有何区别：

skip\_list.c中：

**初始化一个slnode节点：**


```
1 void skiplist_node_init(skiplist_node *node)
2 {
3     memset(node, 0, sizeof(skiplist_node)); /*内存置零，就是准备留给新对象使用*/
4 }
```

**初始化一个跳表：**

- 跳表最大级数是8；
- 初始化next和prev指向节点本身；

```
1 //MaxNumberOfLevels是跳表的级数，这里定义最多8级
2 #define MaxNumberOfLevels 8
3 #define MaxLevel (MaxNumberOfLevels - 1)
4
5 void skiplist_init(skiplist_node *slnode)
6 {
7     int i;
8
9     slnode->key = 0xFFFFFFFFFFFFFFFF;
10    slnode->level = 0;
11    slnode->value = NULL;
12    for (i = 0; i < MaxNumberOfLevels; i++)
13        slnode->next[i] = slnode->prev[i] = slnode; /*初始slnode的next和prev指
14    自己*/
15 }
```

个人理解：这里的MaxNumberOfLevels级数不是指下图中的黄色底标的数字，是指从左到右的箭头的层数。





创建一个新的空表：

```
1 skiplist *new_skiplist(skiplist_node *slnode)
2 {
3     skiplist *l = kzalloc(sizeof(skiplist), GFP_ATOMIC);
4
5     BUG_ON(!l);
6     l->header = slnode;
7     return l;
8 }
```

销毁一张表：

```
1 void free_skiplist(skiplist *l)
2 {
3     skiplist_node *p, *q;
4
5     p = l->header;
6     do {
7         q = p->next[0];
8         p->next[0]->prev[0] = q->prev[0];
9         skiplist_node_init(p);
10        p = q;
11    } while (p != l->header);
12    kfree(l);
13 }
```

插入一个节点：

```
1 void skiplist_insert(skiplist *l, skiplist_node *node, keyType key,
2 valueType value, unsigned int randseed)
3 {
4     skiplist_node *update[MaxNumberOfLevels];
5     skiplist_node *p, *q;
6     int k = l->level;
7
8     /*步骤1，从最高层一层一层往下找，并更新update数组，update数组中保存的是每次降一层时的节点*/
9     p = l->header;
10    do {
11        while (q = p->next[k], q->key <= key)
12            p = q;
13        update[k] = p;
14    } while (--k >= 0);
15
16    ++l->entries;
17    /*步骤2，产生一个随机层数level，如果新生成的层数比跳表的层数大，则设置k为跳表当前level大1的层数，并更新update中k层指向header*/
18    k = randomLevel(randseed); /*需要插入的层*/
```

```

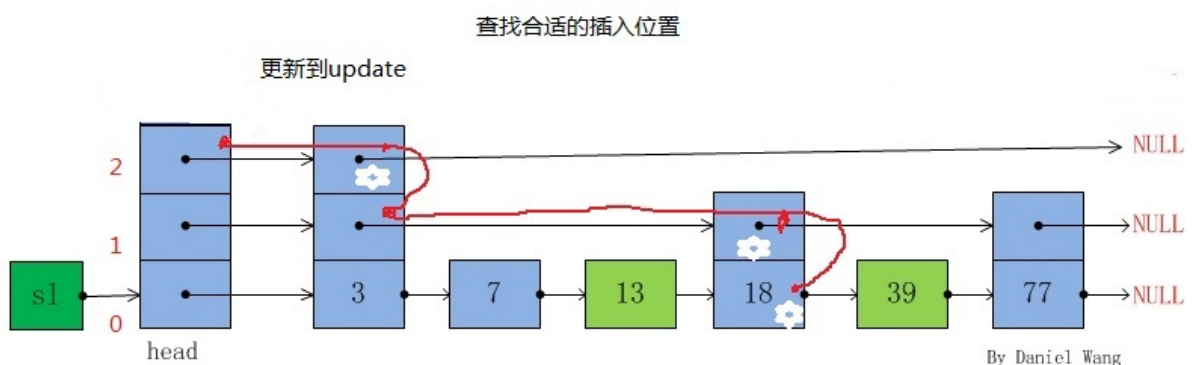
18     if (k > l->level) {
19         k = ++l->level;
20         update[k] = l->header;
21     }
22
23     /*步骤3, 将待插入的节点一层一层的插入*/
24     node->level = k; /*node当中的level记录了该节点从哪一层被插入*/
25     node->key = key;
26     node->value = value;
27     do {
28         p = update[k]; /*这里逐层往下插入到update[k]之后一个元素*/
29         node->next[k] = p->next[k];
30         p->next[k] = node;
31         node->prev[k] = p;
32         node->next[k]->prev[k] = node;
33     } while (--k >= 0);
34 }
35
36 /*步骤2中为何不用担心k超过MaxNumberOfLevels?应该是这个计算随机数时会保证在MaxLevel范围内*/
37 static inline unsigned int randomLevel(const long unsigned int randseed)
38 {
39     return find_first_bit(&randseed, MaxLevel) / 2;
40 }

```

下图中假设要插入的值是25:

### 步骤1

我们需要对于每一层进行遍历并保存这一层中下降的节点(其后继节点为NULL或者后继节点的key大于等于要插入的key), 如下图, 节点中有白色星花标识的节点保存到update数组。



### 步骤2

通过一个随机算法产生一个随机的层数, 但是当这个随机产生的层数level大于当前跳表的最大层数时, 我们此时需要更新当前跳表最大层数到level之间的update内容, 这时应该更新其内容为跳表的头节点head, 想想为什么这么做? 然后就是更新跳表的最大层数。

这么做update[k]=l->header是为啥, 是因为多加了一层时, 只有header么?


对, 我觉得答案就是这样, 当计算随机数k已经大于当前list的最大level了, 则向上取一层, 这一层是新的, 里面没有元素, 所以update[k]需要从header处开始插入

还有为什么会算出来k大于了当前跳表的最大层数? 那样每次都  $k == ++l \rightarrow level$  会不会k超过了最大跳表值?

函数randomLevel中应该限制了MaxLevel (是MaxNumberOfLevel-1)

### 步骤3

根据update[k]开始逐层往下插入节点



删除跳表节点:

```
1 void skiplist_delete(skiplist *l, skiplist_node *node)
2 {
3     int k, m = node->level; /*这里的node-level在插入时设置了*/
4
5     /*for循环从0层开始一直到node所在插入层遍历, 将node节点的next和prev移除*/
6     for (k = 0; k <= m; k++) {
7         node->prev[k]->next[k] = node->next[k];
8         node->next[k]->prev[k] = node->prev[k];
9     }
10    skiplist_node_init(node); /*只是把node节点内存置空, 并未释放这部分内存啊? */
11    /*如果m刚好是最顶层, 删除节点后需要检查下是否m层只剩下header, 如果是删除m层*/
12    if (m == l->level) {
13        while (l->header->next[m] == l->header && l->header->prev[m] == l->header && m > 0)
14            m--;
15        l->level = m;
16    }
17    l->entries--; /*这个entries计数有何用? 待解答*/
18 }
```

## 7.4.2 MuQSS详细设计

### 为何设计MuQSS?

BFS中是所有CPU共享一个runqueue, 这会导致什么呢? 会导致每个CPU都需要去搜索整个runqueue去寻找拥有最早的deadline的进程来调度, 并且不用管该进程原来是哪个CPU调度的, 从而导致BFS的延时会因processed和CPU的数量增加而增加。并且, 单个runqueue会导致CPU之间的锁竞争, 当CPU数量超过16个后, lock contention就很严重了。

MuQSS是BFS的一种进化方案, 改进在哪里?

- 每个CPU都有自己的runqueue
- skiplist跳表取代链表

那么，当初BFS为何只用一个runqueue呢？

是因为有multiple runqueues会需要复杂的交互，因为每个runqueue都只会负责它自己队列的**调度延时和公平性**，这边需要一个复杂的交互系统来保证低延时和公平性，任何增加CPU本地进程调度吞吐量带来优势的同时也会带来劣势，这是因为需要一个复杂的平衡系统，来保证绑定同一个进程到同一个CPU的低延时效果，而不是同一个进程被不同的CPU调度运行。

MuQSS怎么解决多个runqueue带来的劣势问题？

MuQSS通过跳表优先级排序、创新的使用了无锁检查（当它需要因为降低延时需求或者CPU平衡等理由来从其他队列中获取更早的deadline的任务时）。MuQSS仍然没有balancing系统，选择允许下一个任务调度决策和任务唤醒CPU来实现平衡。

## 详细设计

### 1. 定制的skiplist实现

MuQSS使用固定的**8 level**跳表，不是动态分配的，这样使得每个队列仅可将 $O(\log N)$ 扩展为64k个任务（**这个地方没搞懂**），但是呢，每个CPU都有一个runqueue的话，这样 $O(\log N)$ 最多可扩展 $64k * \text{CPUs}$ 个任务，和CPU数量相关了

### 2. 任务插入

MuQSS任务插队就是一个 $O(\log N)$ 的插入skiplist操作。

### 3. Niffies

jiffies是记录系统启动后到现在的时钟中断次数，它取决于系统的时钟频率，比如1000Hz，那么产生时钟中断是没1/1000s一次，也即1ms一次。

niffies和jiffies不同，niffies是一直单调递增的定时器，纳秒单位，Niffies是根据高分辨率TSC计时器针对每个运行队列计算的，并且为了保持公平性，每当两个运行队列同时锁定时，CPU之间就会进行同步

### 4. virtual deadline

虚拟期限？，MuQSS中保证**低延时、调度公平性、优先级**的关键核心机制是**virtual deadline mechanism**。

**rr\_interval: round robin interval**，该参数可通过proc系统调节，作用是：当两个任务具有相同的nice级别时（普通进程SCHED\_NORMAL或者SCHED\_OTHER），该进程能够运行的最大时间；或者换个角度说，两个相同优先级任务的最大延迟时间。

当一个任务需要CPU时间，它被配置的时间片（**time\_slice**）等于一个rr\_interval和一个virtual deadline，（这里如何理解？），virtual deadline如何计算：

```
niffies + (prio_ratio * rr_interval)
```

其中：

- prio\_ratio：优先级，是和nice -20的基线进行比的比率，每增加一个nice level，prio\_ratio增加10%；

- deadline: (deadline调度器是根据deadline来选择调度的, 最先到达截止时间点的进程被有优先调度); 截止时间点, 是个虚拟时间, 用于比较接下来调度运行那个任务。

选择哪个进程该运行, 通常有三种情况:

- 时间片耗尽, 进程会被重新调度, 时间片也会被分配, deadline也会按照上面的公式进行重计算;
- 进程进入睡眠sleep状态, 会让出CPU, 这个过程中, time\_slice时间片和deadline不会改变, 该进程下次被调度时还会恢复;
- 抢占, 一个新的任务比当前正在运行的任务有更高的优先级, 可以抢占。

在前两种情况中, deadline是选择下一个运行任务的关键要素点。

The CPU proportion of different nice tasks works out to be approximately the  
(prio\_ratio difference)^2

The reason it is squared is that a task's deadline does not change while it is running unless it runs out of time\_slice. Thus, even if the time actually passes the deadline of another task that is queued, it will not get CPU time unless the current running task deschedules, and the time "base" (niffies) is constantly moving.

## 5. 任务查找

由于在skiplist中, 任务已经预先根据调度的预期顺序排序了, 通常选择下一个待运行的任务就是选择**0 level的第一个entry入口任务**,

查找的时间复杂度是O(k), 这里的k是CPU个数。

## 6. 延时

通过使用虚拟期限来控制正常任务的调度顺序, 可以确保每个运行队列的队列到激活延迟都受rr\_interval可调参数约束, 该参数默认设置为6ms。这意味着与CPU绑定的任务等待的最长时间将与正在运行的任务的数量成正比, 在通常情况下, 每个CPU 0-2个正在运行的任务, 将低于7ms的阈值(人类能感到抖动的阈值)。

### 7.4.3 源码简析

**rq**

MuQSS对rq结构体进行了定制修改:

```
1  /*
2   * This is the main, per-CPU runqueue data structure.
3   * This data should only be modified by the local cpu.
4   */
5  struct rq {
6      raw_spinlock_t *lock;
7      raw_spinlock_t *orig_lock;
8
9      struct task_struct __rcu *curr;
10     struct task_struct *idle;
11     struct task_struct *stop;
12     struct mm_struct *prev_mm;
13 }
```

```

14     unsigned int nr_running;
15     /*
16      * This is part of a global counter where only the total sum
17      * over all CPUs matters. A task can increase this counter on
18      * one CPU and if it got migrated afterwards it may decrease
19      * it on another CPU. Always updated under the runqueue lock:
20      */
21     unsigned long nr_uninterruptible;
22 #ifdef CONFIG_SMP
23     unsigned int          ttwu_pending;
24 #endif
25     u64 nr_switches;
26
27     /* Stored data about rq->curr to work outside rq lock */
28     u64 rq_deadline;
29     int rq_prio;
30
31     /* Best queued id for use outside lock */
32     u64 best_key;
33
34     unsigned long last_scheduler_tick; /* Last jiffy this RQ ticked */
35     unsigned long last_jiffy; /* Last jiffy this RQ updated rq clock */
36     u64 niffies; /* Last time this RQ updated rq clock */
37     u64 last_niffy; /* Last niffies as updated by local clock */
38     u64 last_jiffy_niffies; /* Niffies @ last_jiffy */
39
40     u64 load_update; /* when we last updated load */
41     unsigned long load_avg; /* Rolling load average */
42 #ifdef CONFIG_HAVE_SCHED_AVG_IRQ
43     u64 irq_load_update; /* when we last updated IRQ load */
44     unsigned long irq_load_avg; /* Rolling IRQ load average */
45 #endif
46 #ifdef CONFIG_SMT_NICE
47     struct mm_struct *rq_mm;
48     int rq_smt_bias; /* Policy/nice level bias across smt siblings */
49 #endif
50     /* Accurate timekeeping data */
51     unsigned long user_ns, nice_ns, irq_ns, softirq_ns, system_ns,
52         iowait_ns, idle_ns;
53     atomic_t nr_iowait;
54
55 #ifdef CONFIG_MEMBARRIER
56     int membarrier_state;
57 #endif
58
59     skiplist_node *node;
60     skiplist *sl;
61 #ifdef CONFIG_SMP
62     struct task_struct *preempt; /* Preempt triggered on this task */
63     struct task_struct *preempting; /* Hint only, what task is preempting
64 */
65
66     int cpu; /* cpu of this runqueue */
67     bool online;
68
69     struct root_domain *rd;
70     struct sched_domain *sd;

```

```

71     unsigned long cpu_capacity_orig;
72
73     int *cpu_locality; /* CPU relative cache distance */
74     struct rq **rq_order; /* shared RQs ordered by relative cache distance
75     */
76
77     struct rq **cpu_order; /* RQs of discrete CPUs ordered by distance */
78
79     bool is_leader;
80     struct rq *smp_leader; /* First physical CPU per node */
81 #ifdef CONFIG_SCHED_THERMAL_PRESSURE
82     struct sched_avg avg_thermal;
83 #endif /* CONFIG_SCHED_THERMAL_PRESSURE */
84 #ifdef CONFIG_SCHED_SMT
85     struct rq *smt_leader; /* First logical CPU in SMT siblings */
86     cpumask_t thread_mask;
87     bool (*siblings_idle)(struct rq *rq);
88     /* See if all smt siblings are idle */
89 #endif /* CONFIG_SCHED_SMT */
90 #ifdef CONFIG_SCHED_MC
91     struct rq *mc_leader; /* First logical CPU in MC siblings */
92     cpumask_t core_mask;
93     bool (*cache_idle)(struct rq *rq);
94     /* See if all cache siblings are idle */
95 #endif /* CONFIG_SCHED_MC */
96 #endif /* CONFIG_SMP */
97
98 #ifdef CONFIG_IRQ_TIME_ACCOUNTING
99     u64 prev_irq_time;
100 #endif /* CONFIG_IRQ_TIME_ACCOUNTING */
101 #ifdef CONFIG_PARAVIRT
102     u64 prev_steal_time;
103 #endif /* CONFIG_PARAVIRT */
104 #ifdef CONFIG_PARAVIRT_TIME_ACCOUNTING
105     u64 prev_steal_time_rq;
106 #endif /* CONFIG_PARAVIRT_TIME_ACCOUNTING */
107
108     u64 clock, old_clock, last_tick;
109     /* Ensure that all clocks are in the same cache line */
110     u64 clock_task ____cacheline_aligned;
111     int dither;
112
113     int iso_ticks;
114     bool iso_refractory;
115
116 #ifdef CONFIG_HIGH_RES_TIMERS
117     struct hrtimer hrtimer;
118 #endif
119
120     int rt_nr_running; /* Number real time tasks running */
121 #ifdef CONFIG_SCHEDSTATS
122     /* latency stats */
123     struct sched_info rq_sched_info;
124     unsigned long long rq_cpu_time;
125     /* could above be rq->cfs_rq.exec_clock + rq->rt_rq.rt_runtime ? */
126
127     /* sys_sched_yield() stats */
128     unsigned int yld_count;

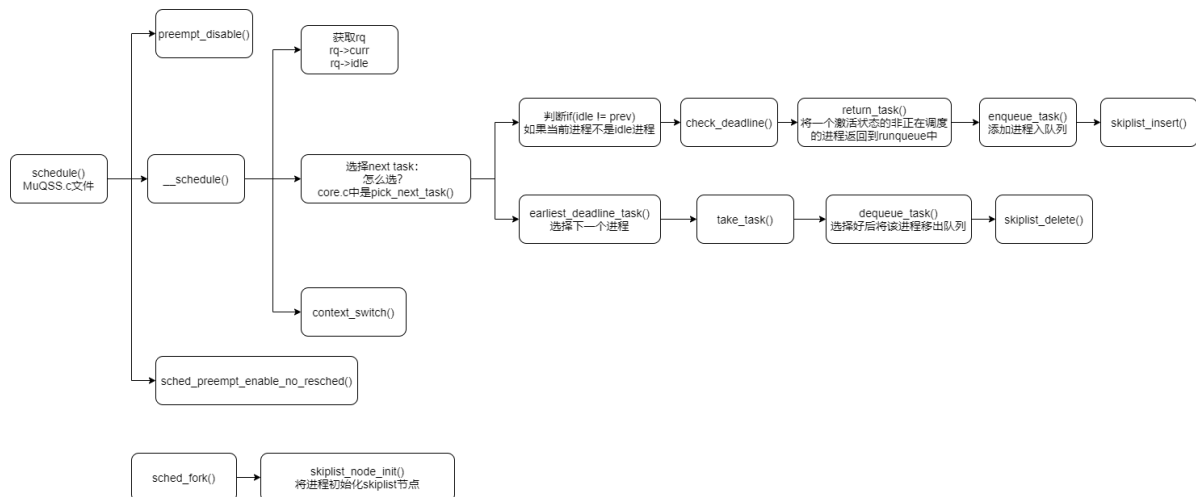
```

```

128
129     /* schedule() stats */
130     unsigned int sched_switch;
131     unsigned int sched_count;
132     unsigned int sched_goidle;
133
134     /* try_to_wake_up() stats */
135     unsigned int ttwu_count;
136     unsigned int ttwu_local;
137 #endif /* CONFIG_SCHEDSTATS */
138
139 #ifdef CONFIG_CPU_IDLE
140     /* Must be inspected within a rcu lock section */
141     struct cpuidle_state *idle_state;
142 #endif
143 }

```

MuQSS.c直接将core.c原本的框架修改了，直接不涉及调度类等，集成在MuQSS.c中



## 八、实时调度器

## 九、自定义调度器



## 9.1 框架

框架需不需要修改？

MuQSS中的MuQSS.c对应主线的core.c，其中实现了大部分原主线上的函数比如context\_switch()等等；

**TODO:**

实现一个简单的先进先出的调度器，框架没必要修改，还是用原来的core.c中实现的即可；

## 9.2 调度类

现有的主线core.c中涉及调度类stop > deadline > rt > fair > idle，共5个调度类；

MuQSS暂未找到调度类的用法，其很多功能都是在MuQSS.c封装实现，比如：

enqueue\_task(), dequeue\_task(), yield\_to() 等功能函数，主线core.c中是封装后通过class->enqueue\_task()钩子函数来调用具体的调度器算法的功能实现，MuQSS中是enqueue\_task()直接实现；

**TODO:**

设计一个sched\_class，集成到内核中；

## 9.3 具体实现

### 9.3.1 数据结构

**TODO:**

使用双链表或者跳表

runqueue需要重新设计一个rq，并嵌入到struct rq{}中去

runqueue

sched\_class

policy

9.3.2 调度算法

TODO:

先进先出的算法，不涉及抢占，时间片用完或者自动退出才进行进程调度切换；

9.3.3 函数功能

TODO:

根据sched\_class中的钩子函数，参照fair.c和MuQSS，实现每个钩子函数即可；

9.4 功能测试验证

9.4.1 编译

TODO:

涉及内核调度器修改，编译恐怕容易出错，跑起来容易导致panic

9.4.2 验证

TODO:

- 将其他的sched\_class全部阉割掉，只留自开发的FIFO调度器；
- FIFO调度器和主线上调度器并存，用户设置进程使用哪种调度器比如FIFO来验证性能；


十、调度器框架

模仿fair\_sched\_class来修改编码fifolist\_sched\_class：

core.c中涉及的fair\_sched\_class有以下：

sched_init()	初始化调度器函数，由start_kernel调用
__setscheduler()	sched_setscheduler()函数会调用，用于修改线程的sched policy或者priority
rt_mutex_setprio()	设置当前任务的优先级
set_load_weight()	
set_task_cpu()	
sched_fork()	fork的进程调度
pick_next_task	选择下一个调度的进程任务

## 10.1 调度器初始化



core.c中的sched\_init()函数到底初始化了什么？

```
1 {
2     unsigned long ptr = 0;
3     int i;
4
5     /* Make sure the linker didn't screw up */
6     BUG_ON(&idle_sched_class + 1 != &fair_sched_class ||
7           &fair_sched_class + 1 != &rt_sched_class ||
8           &rt_sched_class + 1 != &dl_sched_class);
9     #ifdef CONFIG_SMP
10        BUG_ON(&dl_sched_class + 1 != &stop_sched_class);
11    #endif
12
13    wait_bit_init();
14
```

```

15 #ifdef CONFIG_FAIR_GROUP_SCHED
16     ptr += 2 * nr_cpu_ids * sizeof(void **);
17 #endif
18 #ifdef CONFIG_RT_GROUP_SCHED
19     ptr += 2 * nr_cpu_ids * sizeof(void **);
20 #endif
21     if (ptr) {
22         ptr = (unsigned long)kzalloc(ptr, GFP_NOWAIT);
23
24 #ifdef CONFIG_FAIR_GROUP_SCHED
25         root_task_group.se = (struct sched_entity **)ptr;
26         ptr += nr_cpu_ids * sizeof(void **);
27
28         root_task_group.cfs_rq = (struct cfs_rq **)ptr;
29         ptr += nr_cpu_ids * sizeof(void **);
30
31         root_task_group.shares = ROOT_TASK_GROUP_LOAD;
32         init_cfs_bandwidth(&root_task_group.cfs_bandwidth);
33 #endif /* CONFIG_FAIR_GROUP_SCHED */
34 #ifdef CONFIG_RT_GROUP_SCHED
35         root_task_group.rt_se = (struct sched_rt_entity **)ptr;
36         ptr += nr_cpu_ids * sizeof(void **);
37
38         root_task_group.rt_rq = (struct rt_rq **)ptr;
39         ptr += nr_cpu_ids * sizeof(void **);
40
41 #endif /* CONFIG_RT_GROUP_SCHED */
42     }
43 #ifdef CONFIG_CPUMASK_OFFSTACK
44     for_each_possible_cpu(i) {
45         per_cpu(load_balance_mask, i) = (cpumask_var_t)kzalloc_node(
46             cpumask_size(), GFP_KERNEL, cpu_to_node(i));
47         per_cpu(select_idle_mask, i) = (cpumask_var_t)kzalloc_node(
48             cpumask_size(), GFP_KERNEL, cpu_to_node(i));
49     }
50 #endif /* CONFIG_CPUMASK_OFFSTACK */
51
52     init_rt_bandwidth(&def_rt_bandwidth, global_rt_period(),
53         global_rt_runtime());
54     init_dl_bandwidth(&def_dl_bandwidth, global_rt_period(),
55         global_rt_runtime());
56
57 #ifdef CONFIG_SMP
58     init_defrootdomain();
59 #endif
60
61 #ifdef CONFIG_RT_GROUP_SCHED
62     init_rt_bandwidth(&root_task_group.rt_bandwidth,
63         global_rt_period(), global_rt_runtime());
64 #endif /* CONFIG_RT_GROUP_SCHED */
65
66 #ifdef CONFIG_CGROUP_SCHED
67     task_group_cache = KMEM_CACHE(task_group, 0);
68
69     list_add(&root_task_group.list, &task_groups);
70     INIT_LIST_HEAD(&root_task_group.children);
71     INIT_LIST_HEAD(&root_task_group.siblings);
72     autogroup_init(&init_task);

```

```

71 #endif /* CONFIG_CGROUP_SCHED */
72
73     for_each_possible_cpu(i) {
74         struct rq *rq;
75
76         rq = cpu_rq(i);
77         raw_spin_lock_init(&rq->lock);
78         rq->nr_running = 0;
79         rq->calc_load_active = 0;
80         rq->calc_load_update = jiffies + LOAD_FREQ;
81         init_cfs_rq(&rq->cfs);
82         init_rt_rq(&rq->rt);
83         init_dl_rq(&rq->dl);
84 #ifdef CONFIG_FAIR_GROUP_SCHED
85         INIT_LIST_HEAD(&rq->leaf_cfs_rq_list);
86         rq->tmp_alone_branch = &rq->leaf_cfs_rq_list;
87         /*
88          * How much CPU bandwidth does root_task_group get?
89          *
90          * In case of task-groups formed thr' the cgroup filesystem, it
91          * gets 100% of the CPU resources in the system. This overall
92          * system CPU resource is divided among the tasks of
93          * root_task_group and its child task-groups in a fair manner,
94          * based on each entity's (task or task-group's) weight
95          * (se->load.weight).
96          *
97          * In other words, if root_task_group has 10 tasks of weight
98          * 1024 and two child groups A0 and A1 (of weight 1024 each),
99          * then A0's share of the CPU resource is:
100         *
101         *   A0's bandwidth = 1024 / (10*1024 + 1024 + 1024) = 8.33%
102         *
103         * We achieve this by letting root_task_group's tasks sit
104         * directly in rq->cfs (i.e root_task_group->se[] = NULL).
105         */
106         init_tg_cfs_entry(&root_task_group, &rq->cfs, NULL, i, NULL);
107 #endif /* CONFIG_FAIR_GROUP_SCHED */
108
109         rq->rt.rt_runtime = def_rt_bandwidth.rt_runtime;
110 #ifdef CONFIG_RT_GROUP_SCHED
111         init_tg_rt_entry(&root_task_group, &rq->rt, NULL, i, NULL);
112 #endif
113 #ifdef CONFIG_SMP
114         rq->sd = NULL;
115         rq->rd = NULL;
116         rq->cpu_capacity = rq->cpu_capacity_orig = SCHED_CAPACITY_SCALE;
117         rq->balance_callback = NULL;
118         rq->active_balance = 0;
119         rq->next_balance = jiffies;
120         rq->push_cpu = 0;
121         rq->cpu = i;
122         rq->online = 0;
123         rq->idle_stamp = 0;
124         rq->avg_idle = 2*sysctl_sched_migration_cost;
125         rq->max_idle_balance_cost = sysctl_sched_migration_cost;
126
127         INIT_LIST_HEAD(&rq->cfs_tasks);
128

```

```

129     rq_attach_root(rq, &def_root_domain);
130 #ifdef CONFIG_NO_HZ_COMMON
131     rq->last_blocked_load_update_tick = jiffies;
132     atomic_set(&rq->nohz_flags, 0);
133
134     rq_csd_init(rq, &rq->nohz_csd, nohz_csd_func);
135 #endif
136 #endif /* CONFIG_SMP */
137     hrtick_rq_init(rq);
138     atomic_set(&rq->nr_iowait, 0);
139 }
140
141     set_load_weight(&init_task, false);
142
143     /*
144      * The boot idle thread does lazy MMU switching as well:
145      */
146     mmgrab(&init_mm);
147     enter_lazy_tlb(&init_mm, current);
148
149     /*
150      * Make us the idle thread. Technically, schedule() should not be
151      * called from this thread, however somewhere below it might be,
152      * but because we are the idle thread, we just pick up running again
153      * when this runqueue becomes "idle".
154      */
155     init_idle(current, smp_processor_id());
156
157     calc_load_update = jiffies + LOAD_FREQ;
158
159 #ifdef CONFIG_SMP
160     idle_thread_set_boot_cpu();
161 #endif
162     init_sched_fair_class();
163
164     init_schedstats();
165
166     psi_init();
167
168     init_uclamp();
169
170     scheduler_running = 1;
171 }

```

